

2017-2 Distributed Systems Part2

45748 characters in 7469 words on 1242 lines

Florian Moser

February 20, 2018

1 Introduction

reasons for distributed systems

geography (large firms are active worldwide)
parallelism (multicore to speed up computation)
reliability (replication to prevent data loss)
availability (replication for fast access)

problems of distribution

coordination (consistency, agreement, consensus)
probability high that some machine in cluster is down

2 Fault-Tolerance & Paxos

2.1 definitions

node

single actor in system
total amount of nodes denoted by n

client / server

client node wants to manipulate data on server node

sending one at a time

only ever sends next message if previous message ACKed

sequence number

unique number attached to every message
allows to discover duplicates

state replication

archived by set of nodes which execute commands in same order
fundamental property

ticket

weaker form of lock, reissuable, expires if new one arrived
reissuable to deal with crashes, implement with counter

2.2 models

message passing model MPM

set of nodes, each performs computations, sends messages to all others

MPM with message loss MPMML

any message may be lost on the way to receiver
message corruption better than total loss (can use checksums)

variable message delay VMD

different transaction times for each message possible

2.3 algorithms

2.3.1 naive client-server

client sends commands to server

remarks

not robust against message loss

2.3.2 client-server with ACK

client sends commands one at a time to server
servers responds with ACK
client resends if no ACK received

remarks

include sequence numbers to avoid double execution
basis of TCP and other reliable protocols

inconsistent state possible with multiple servers

due to reordering because of VMD
proof with $x+1$, $x*2$ received at different orders

2.3.3 state replication with serializer

client sends commands one at a time to serializer

serializer forwards one at a time to servers

serializer notifies client of success

remarks

also called master-slave
serializer is single point of failure

2.3.4 two-phase protocol

(phase 1)

client asks servers for lock

(phase 2)

if (client has locks from all) then send command, return locks
else give locks back, wait and restart phase 1

remarks

applications include 2PC, 2PL, 3PL
needs to have all servers available to function

2.3.5 naive ticket protocol

(phase 1)

client asks servers for ticket

(phase 2)

if (majority replied) then

client sends command with ticket number

servers respond if ticket number is max

(phase 3) continue if majority replied

if (majority replied) then

client tells servers to execute

else

client waits and restarts phase 1

remarks

if client slow with execute, other may slip in own command

2.3.6 paxos

(init)

client has command c , ticket number $t=0$

server has $T_{max} = 0$, $T_c = 0$, $C = \text{null}$

(phase 1)

client increases t , sends (t, c) to server

if $(t > T_{max})$ then server updates T_{max} , responds with (T_{store}, C)

(phase 2) continue if majority replied

set $c = \max T_{store} C??$ c , send propose(t, c) to same

if $(t == T_{max})$ then server sets C , T_{store} , answers success

(phase 3) continue if majority replied

sends execute(c) to all

remarks

clients can abort and restart at any point

randomize waiting times and send NACK for better performance

needs majority of servers up, needs trusted clients

instance of paxos decides on single command

run paxos in parallel to decide on multiple commands

may not terminate

worst-case

all clients start at same time, same timeout, same initial ticket

then they keep invalidating each other at phase 2, can't enter phase 3

improve

do not use different initial ticket numbers, leads to starvation

use exponential backoff for timeouts in phase 2, phase 3

proposal chosen is kept

chose if majority of servers accept

only can happen if majority replied in (phase 1)

therefore for all t only single c can be chosen

3 Consensus

3.1 definitions

consensus

correct nodes must decide on single value from the proposed values
agreement (1) & termination (2) & validity (3) (any-input)
nodes have no shared memory, focus on algorithms with progress

asynchronous runtime

time units (delay of single message) from start to end at worst case
cannot make assumptions about maximum delay in algorithm

configuration

fully defined system at specific point in time
includes all messages in transit & state of all nodes
initial configuration is C_0 , all nodes have sent first message

univalent configuration

if decision value fixed no matter of what happens afterwards

v-valent configuration

univalent configuration for value v

bivalent

if decision value not fixed (yet)

transition T

from C_1 to C_2 , is event (u, m) with node u receiving message m
in C_2 , m arrived, u changed state, new messages from u in transit

configuration tree

directed tree of C s, top is C_0 , edges are all possible T s
every algorithm has one tree per selection of input values
leaves are univalent end states
every possible path to every leave is one possible execution
if node u crashes in C remove all $(u, *)$ transitions

critical configuration C

if C is bivalent but all C 's below are univalent

3.2 models

asynchronous model

algorithms are event based (on receive, do ...)
no synchronized time
messages are received in finite but unbound time

3.3 proofs

there are bivalent C_0

build a = array of initial values, index of a determines number of 0
 a_0 is 0-valent, a_n is 1-valent, there must be turning point a_i
node i crashes, the other nodes must terminate & decide on value

(u_1, m_1) (u_2, m_2) end in same C

as consume/produce same messages, states, they end in same C

system must reach critical configuration to terminate

assume bivalent start, always progress in bivalent configuration

no termination if crash at critical configuration

define C with t_0 0-valent and t_1 1-valent
must happen on same node; crash that node and remove transitions

if $f > 0$ & deterministic algorithm then no termination

crash node with critical configuration

$f \geq n/2$ can't archive consensus

define S_1 all nodes are 1, S_2 half 0/half 1, S_3 all are 0
build set N and N' with each $n/2$ nodes, crash N'
if N is all 1, can't know if in S_1 or S_2
this proof sketch is useful for similar problems

consensus with f edge failures

assuming fully connected network, $n*(n-1)/2$ edges
largest f is path between remaining nodes
smallest f partitions, $n-1$

3.4 algorithms

3.4.1 sending ACK after ACK

receiver sends ACK after receiving message
sender confirms ACK with ACK

remarks

can never terminate, need to continue indefinitely
proof by assuming an algorithm exists, and then losing its last ACK

3.4.2 naive consensus

node broadcasts its value, waits for all others
chooses minima

remarks

does not tolerate crashes

3.4.3 randomized consensus (async)

(init)

client has $v = \{0,1\}$, round = 1, decided = false
broadcast $myValue(v, round)$
(propose) when majority of $myValue$ messages received
if (all same value) then propose($v, round$)
else propose($null, round$)
if (decided) $myValue(v, round+1)$ and terminate
(adapt) when majority of propose messages received
if (all same proposed) $v=proposed$ and decided=true
elseif (at least one proposed) $v=proposed$
else choose v randomly
round += 1, $myValue(v, round)$
(restart round at propose)

remarks

progress ensured at (propose) and (adapt) for $f < n/2$
cannot set v deterministic, proof using a partition

validity fulfilled

trivial for univalent input, else result does not matter

agreement fulfilled

terminates if decided is true, is true if all propose same value
send of propose possible if majority same $myValue v$ received
it follows all have received at least one propose with $myValue v$
nodes which can't decide adapt own value and decide next round

termination fulfilled

all pick same value with probability 2^{-n}
constant messages per round, therefore $O(2^n)$

3.4.4 shared coin (async)

broadcast coin $c = 0$ with $1/n$, else $c = 1$
wait for $n-f$ coins
put all coins in set and broadcast
wait for $n-f$ sets
if (0-coin in any sets) decide 0 else decide 1

remarks

at most f nodes crash therefore the two waits are OK
all coins are seen by all correct nodes

3.4.5 exercise algorithms

bandwidth limitation (sync)

assume nodes with ID, no crashes, single message per round
determine leader by $id = 1$, leader sends value to $id = 2$
generalize, only $\log(n)$ rounds necessary

consensus in a grid

each node broadcast own value, resends all received values
waits till no new information received
 $l+1$ rounds for $l = w+h$ of grid
need only single byzantine node to deceive corner node

4 Byzantine Agreement

4.1 definitions

byzantine

node with arbitrary behaviour, also includes malicious

byzantine agreement

finding consensus with byzantine nodes in network
agreement (1), termination (2), any-input validity (3)

f-resilient

system still operative with f byzantine nodes

any-input validity ("normal")

decision value is input of any node

correct-input validity

decision value is input of a correct node

all-same validity

if all correct nodes propose same value, this will be decided

median validity

decision value is value close to the median of correct nodes

synchronous runtime

number of rounds from start to end of worst case

bounds

upper bound if problem can be solved in that time

lower bound if problem can't be solved in less time

tight bound if lower=upper bound

4.2 models

synchronous model

operation in rounds, each round messages are sent and received

4.3 proofs

all-same validity & byzantine agreement needs $n > 3$

for $n=3$ can't differentiate friends from fiend

byzantine agreement needs $f < n/3$

combine nodes in 3 supernodes, then use proof for $n > 3$

need at least $f+1$ rounds to decide for minima

f times contact single neighbour which crashes afterwards

general proof also possible

paxos fails

assume two majorities which overlap in one server

this server can decide for different values in the two majorities

4.4 algorithms

4.4.1 byzantine agreement (for $f=1$)

node u has value x

(round 1)

send tuple (u, x) to all others, receive tuples from all others

store received tuples in set S

(round 2)

send set S to all others, receive sets U from all others

choose smallest value from tuples contained > 1 in $(U + S)$

remarks

in round 2 only received tuples are resend, but not own

all correct nodes will have same U

archive all-same validity with multiple occurrence of min-value

4.4.2 king algorithm (for $f < n/3$)

node has value x

for phase=1 till $f+1$

(round 1)

value(x)

(round 2)

if $(\#value(v) > n-f)$ then propose(v)

if $(\#propose(z) > f)$ then $x=z$

(round 3)

predefined king of phase broadcasts his value(w)

if $(\#propose(x) < n-f)$ then $x=w$

endfor

remarks

first adaptation needed for king

second adaptation for backups

after a correct node had been king, no changes any more

correct nodes propose only one value

at least one phase has a correct king

additionally all-same validity

4.4.3 asynchronous byzantine agreement ($f < n/10$)

node has value x , $r = 1$, decided = false

propose(x , r)

(loop) wait for $n-f$ propose message of current round

if $(\#propose y > n/2 + 3f + 1)$ then propose(y , $r+1$), terminate

if $(\#propose y > n/2 + f + 1)$ then $x = y$

else choose x randomly

$r+=1$, propose (x , r)

restart (loop)

remarks

$x=y$ only happens for single y at correct nodes

all-same validity

when than $n-2f$ propose from correct nodes received

then $n-2f > (n/2 + 3f + 1)$

agreement

when $(n/2 + 3f + 1)$ was fulfilled, then $-2f$ worst case

but will still adopt y , and then terminate next round

termination

with probability $1/2^{(n-f)+1}$

5 Authenticated Agreement

5.1 definitions

signature

nodes can sign a message to verify origin

5.2 models

system model SM

$n = 3f + 1$, unbound number of clients which send requests

messages are asynchronous, have variable delay, can get lost

5.3 proofs

quorum intersection $2f + 1$ in SM

intersection of two $2f+1$ sets has at least one correct node

5.4 algorithms

5.4.1 byzantine agreement with authentication (sync)

primary has input, all nodes can sign messages with signature

(primary p)

if (input) send value(1)_p

decide input and terminate

(secondary u)

for $i=0$ till f

if $(\#received\ messages > i\ and\ value(1)_p\ contained)$ then

broadcast all received messages + value(1)_u, terminate

remarks

solves for any number of failures, signatures help detect byzantine

byzantine primary

to avoid a byzantine primary to dictate run the algorithm in parallel

need $2f+1$ nodes ($f < n/2$) as primary, choose where $\#result > f+1$

more than 0-1

primary always broadcast

secondary checks that only single value received from primary

5.4.2 practical byzantine fault tolerance (async)

view

integer which determines configuration state of protocol

primary, backups

for view v , node $u = v \bmod n$ is the primary, others are backups

accepted messages

authenticated (signed), protocol-correct, same view messages

faulty timer

started while waiting for response from primary

trigger view change if timeout occurs

prepared certificate PC

$2f+1$ prepare messages (can include own)

new-view certificate NVC

$2f+1$ view_change messages (can include own)

remarks

signatures verify sender

in each view, there is always a primary, the others are backups

if primary is byzantine, backups can initiate a view change

correct primary chooses dense sequence numbers sn

if one node executed (request, sn) eventually all other will too

nodes collect $2f+1$ confirmation messages before executing (r , sn)

if client receives $f+1$ replies it can assume execution

if prepare certificate obtained, no others exists

request accept

(phase 1, primary p , view v , sequence number s)

accepts request from client

send pre-prepare(v , s , r , p)- p

(phase 1, backup b)

accepts request from client and relays to primary

remarks request accept

client sends request to all servers

primary sends pre-prepare else byzantine could force view_changes
(by sending distinct pre-prepare to all nodes, faulty timers expire)

prepare & pre-prepare

(phase 2, backup b)

accept pre-prepare(v, s, r, p)_p

verify p is primary of v, verify first pre-prepare for (v,s)

send prepare(v, s, r, b)_b

(phase 3, node i)

wait for 2f prepare matching (v,s,r) (PC)

send commit(v,s,i)

wait for 2f+1 commit matching (v,s)

execute r after all lower r's have been executed

send reply to client

view change initialize

(after faulty timer at backup b has expired)

stops accepting for v

send view_change(v+1, P.i, i) with P.i are PC's already established

view change protocol

(new primary p of view v+1)

wait for 2f+1 view_change messages, put in V

O is set of all pre-prepare with PC from V, adapted for v+1

O has pre-prepare(v+1, s, null, p)_p for all s with missing PC

send new_view(v+1, V, O, p)_p to all nodes

start processing with s_max + 1 from O

(backup b of view v+1)

accept new-view(v+1, V, O, p)_p, stops accepting for v

verify p is primary, verify O correctly constructed from V

if (verify OK) then start at (phase 2) else trigger view change v+2

remarks

unique sequence numbers even across views, so each (v,r) is unique

after view_change message, faulty timer if new primary is byzantine

6 Quorum Systems

6.1 definitions

quorum

quorum is a subset of all nodes, such that any two overlap

minimal if its the smallest subset possible

access strategy

probability for each node that it is accessed

work of quorum

number of nodes in one quorum

work of access strategy

expected number of nodes accessed

work of quorum system W(Q)

work of best access strategy

sumof (p_quorum * quorum_size) for all quorums

load of node

probability that it is accessed

sumof (p_quorum) for all quorums node is part of

load of access strategy

maximum load of any node from the system

load of quorum system L(Q)

load of best access strategy

work vs load

work is what has to be done, load is how well its distributed

work is usually a big number (3), load is small (1/3)

uniform access strategy

work is amount of nodes per quorum

load is work(Q)/#quorums

failure probability F_p

p a quorum systems fails assuming nodes fail with fixed p

asymptotic failure probability afp

p when $n \rightarrow \infty$

f-resilience R(Q)

if f nodes can fail but quorum still possible

f-disseminating

assumes self-verifying messages

(1) intersection of two quorum systems contains f+1 nodes

(2) for f byzantine nodes there is quorum system without one

cannot double-spend information (1), cannot simply crash (2)

need more than 3f nodes

$L \geq \sqrt{f+1 / n}$ because f+1 element accessed for quorum

f-masking

extension of f-disseminating because can falsify information

(1) intersection of two quorum systems contains 2f+1 nodes

(2) for f byzantine nodes there is quorum system without one

cannot falsify information (1), cannot simply crash (2)

need more than 4f nodes

$L \geq \sqrt{2f+1 / n}$ because 2f+1 elements accessed for quorum

f-opaque quorum system

ensures each quorum accesses more up-to date nodes than others

$n > 5f$, $L(S) \geq 0.5$

s-uniform

if every quorum has exactly s elements

balanced access strategy

if load constant on all nodes of quorum system

6.2 proofs

$L \geq \sqrt{1/n}$

need to access a node in all minimal quorums

6.3 access strategies

6.3.1 primary copy

single node locked

remarks

$W = 1$, $L = 1$, $R = 0$, $afp = 1 - p$

good choice if failure probability is over 1/2

6.3.2 majority system

more than half of the nodes locked

remarks

$W > n/2$, $L > 1/2$, $R < n/2$, $afp = 0$

6.3.3 basic grid quorum system

assume \sqrt{n} element_of natural_numbers

quorum i consists of row & column i

remarks

$W = 2\sqrt{n} - 1$, $L = 1/\sqrt{n}$, $R = \sqrt{n}-1$, $afp = 1$

each quorum has two intersections

totally order nodes

sequential locking strategy

lock nodes ordered by identifiers

release all locks if some already locked

concurrent locking strategy

priority to quorum which already has highest node locked

better systems

for example T form, size can be reduced till \sqrt{n}

6.3.4 B-grid quorum system

$n = d \cdot h \cdot r$ where $d = \# \text{columns}$, $h = \# \text{band}$, $r = \# \text{rows in band}$

quorum has column in each band ($h \cdot r$)

additionally has one element from all columns from single band

quorum of size $d + h \cdot r - 1$

remarks

$W = d + hr - 1$, $L = (d + hr - 1) / n$, $R = O(\sqrt{n})$, $afp = 0$

number of different quorums is $d \cdot n \cdot h \cdot r^{d-1}$

6.3.5 f-masking grid quorum system

each quorum contains one column + (f+1) rows of nodes

for $2f + 1 \leq \sqrt{n}$

remarks

like multiple T over each others

f-masking, hits lower bound

6.3.6 M-Grid quorum system

each column / row has $\sqrt{f+1}$ rows

remarks
L = $\sqrt{f/n}$
f-masking

7 Eventual Consistency & Bitcoin

7.1 definitions

consistency

all nodes agree on current state in system

availability

system is operational and instantly processes requests

partition tolerance

ability to continue operation even while partition exists

quiescent state

no more messages exchanged, shared state consistent

eventual consistency

form of weak agreement, nodes may disagree temporarily but eventually the system reaches quiescent state

7.2 proofs

CAP theorem

assume two nodes sharing state, cannot communicate
update local state (availability) or not (consistency)

7.3 bitcoin

7.3.1 definitions

bitcoin network

randomly connected overlay network
end users run light, not fully validating, version of nodes
tracks funds of address collaboratively

state of bitcoin

unspent transaction outputs (the UTXO) + global parameters
every node holds complete replica of that state

address

hash(public key), network identifier byte, checksum
stored as base58 which avoids ambiguous characters
20bytes long addresses, impractical to brute force
funds associated is sum of all unspent outputs

output

tuple (amount, spending condition)

spending condition

script, cryptographic puzzle, often singed
can be spend or unspent

input

tuple (reference to output, arguments for spending condition)
reference is tuple (h,i)
h is hash of transaction which created output, i specifies index
the arguments solve the spending condition puzzle

transaction

describes transfer of bitcoins, has inputs / outputs
references outputs are removed, outputs added to the UTXO
maybe less output than input, remainder is called fee
added to memory pool with status unconfirmed
after inclusion into a block remove from pool, set status = confirmed
always spend full amount of coins, add output to self for change
(done because it makes agreeing on transactions easier)

doublespent

multiple transactions spent same output, only one is accepted
nodes do not forward conflicting transactions

bitcoin doubling

as bitcoins traceable, those involved in crime need to be washed
criminal trades his traced bitcoin for less, but clean, bitcoins

proof of work

prove to another party certain amount of resources spent
 $\text{SHA}_{256}(\text{SHA}_{256}(c|x)) < 2^{244}/d$ in bitcoin network
difficulty adjusted to 10 minutes finding a new block
gives the network time to synchronize

proof of work function

(1) $F_d(c,x)$ is fast to compute with given d,c,x
(2) for fixed d,c finding x is difficult but feasible

genesis block

first, initial block (maybe hardcoded), root of tree

block

contains transactions, reference to previous block, nonce
is broadcast as soon as valid nonce is found
finder of the block imposes chosen transactions to others

reward transaction

first transaction in a block, has dummy input
sum of outputs is fixed subsidy plus all fees

block chain

longest path from genesis block to leaf
only transaction in this chain are valid

monotonic read consistency

if node sees new value, newer reads by same node will always return it

monotonic write consistency

write operations of same node are serialized, executed in order

read-your-write consistency

if node updates value, newer reads by same node will return it

casual relation

causally related operations include

$w_a(x) \rightarrow w_a(y)$

$r_a(x) \rightarrow w_a(x)$

$w_b(x) \rightarrow r_a(x)$

any transitive combinations

casual consistency

if any casual related operations are seen in same order

7.4 smart contract

7.4.1 definitions

smart contract

agreement between two or more parties
blockchain guarantees correct execution (conflict mediator)

timelocked transactions

define earliest time to be included in chain
only released into network after timelock expired

singlesig / multisig transactions

amount of signatures required to claim output

7.5 algorithms

7.5.1 naive ATM

ATM make withdrawal request to bank
waits response from bank
if (OK) dispense else display error

remarks

connection problem may blocks request

7.5.2 partition tolerant ATM

if (bank reachable) then
sync local view with bank view
display error if user balance insufficient, abort
endif
dispense cash, write logs for synchronization

remarks

partition tolerant, no longer consistency guaranteed

7.5.3 node receives new block

add new node to own tree
if (height increased) then
compute UTXO for path until max_node
cleanup memory pool
endif

remarks

switching paths may results in unconfirmed transactions
smart data structures avoid having to recompute everything

7.5.4 parties create 2by2 multisig output

B sends list with inputs to A, A selects own inputs
A creates transaction $t_1([I_a, I_b], o = c_a + c_b \rightarrow (A, B))$
A creates timelocked $t_2([o], [c_a \rightarrow A, c_b \rightarrow B])$
A signs t_2 and sends it with t_1 to B

B signs both t_i and sends them back to A
A signs t_1 and broadcasts

remarks

s_1 called setup transaction
 s_2 called refund transaction, ensures funds returned
both must be signed by both parties to be valid

7.5.5 micropayment channel, $S \rightarrow R$, capacity c

create 2by2 multisig ($A=R$, $B=S$), $c_a + c_b = c$
 S resigns $t_2([o], [c_r \rightarrow R, c_s \rightarrow S])$ to buy stuff
at end of period, R publishes last t_2 received from S

remarks

$c_r + c_s = c$, reduce c_s with amount to pay
only pay fees once, instantly finalized

7.5.6 unlimited micropayment channel

create micropayment channel
but introduce kickoff transaction t_k between t_1 and t_2
 t_2 only valid if t_k released into blockchain
 t_k needs to be signed at start like t_1
either party can release t_k if wants t_2 to be spend

8 distributed storage

8.1 definitions

diameter

longest distance between any two nodes (using shortest possible)

easy routing

if node does not know result, it should know who to ask

churn

nodes leaving & joining network in time interval (low means little)

topology properties

- (1) homogeneous, no single point of failure, dominant nodes
- (2) ID assigned from range $[1,0)$, can use decimals
- (3) small degree, if possible polyalgorithmic in n
- (4) small diameter, easy, predictable & fast routing

8.2 network topologies

fat tree

like a tree, connection capacity equals leave count
only one possible routing path

mesh $M(m,d)$

like a grid, connects node in x & y direction with others
 $M(m,1)$ is a path, $M(2,2)$ a grid with four nodes
routing simple, only flip one bit at a time

torus $T(m,d)$

like a mesh, but additionally wrap around
 $T(2,3)$ is a die, $T(3,3)$ a cube with 9 nodes each side

butterfly $BF(d)$

d denotes level
constant node degree $(d+1)$, $(2(d+1))^d$ nodes
level 0 is just a point
level 1 adds new line, left/right connected + cross
level 2 adds new line, left/right connected + middle

Cube-Connected-Cycles $CCC(3)$

degree of 3, replaces corners of hypercube with circles
addressing like (a, b) where b cycle position, a corner position

Shuffle-Exchange $SE(d)$

d denotes number of bits
connect if differ in last bit
connect if obtained by cyclic right/left shift

skip list

linked list with additional forward links
operations take $\log n$ expected time
if inserted promote node if no neighbours are promoted

pancake graph

ID is permutation of $1, \dots, d$
nodes connected if first i numbers of ID are flipped
1234 is neighbour of 3214, 2413 is neighbour of 4213

small world graphs

small diameter but large degree nodes (social networks)

expander graph

sparse graph with good connectivity (clusters are connected)
low degree, small diameter, but hard to route

distributed hash table DHT

implements key-value storage, supporting search, insert, delete
can be implemented with hypercube, using first bits of hash
works with butterfly too, use $d+1$ layer as replication
to setup, new joining nodes ask authority, use static IP's
recursive lookup builds path to node with object
good for caching, but request amplification & source hidden
iterative lookup builds direct connection to node with object
more expensive logic, but less load on network

set of hash functions

can reuse the same hash function if set needed
repeated hashing hashes k time for the k -th hash function
salted hashing includes the number k into the message to be hashed

8.3 proofs

diameter must be bigger than $\log(n) / \log(d+1) - 2$

at most d nodes can be reached, as often as diameter

8.4 algorithms

8.4.1 consistent hashing

hash name of file, hash IP/port of node
store file at closest node determined by hash

remarks

could also store only pointers at designated nodes

8.4.2 DHT

hypercube network with $\log(n)$ hypernodes
each hypernode has $\log(n)$ nodes
nodes connected to all others in hypernode
nodes connected to core nodes in other hypernodes
some nodes have to change hypernode to balance
if n shrinks/grows to threshold, hypercube is resized

remarks

assume bounded # of join/leaves occur in worst case manner
assume attacker can crash designated nodes at any point
system is never fully repaired, but always fully functional
at any interval, attacker can crash at most $\log(n)$ nodes

remarks

each node has $\log(\log(n))$ neighbours
nodes are either in core or in periphery of hypernode
 $\log(n)$ for search/insert, $\log(n)$ neighbours with cheating
handles $\log(n)$ churn, but not byzantine, privacy, ...

9 Game Theory

9.1 definitions

game

two players, at least two strategies
every possible outcome (strategy profile) has payoff

social optimum (pareto optimal)

if it maximizes sum of payoffs

dominant strategy

if player is never worse by choosing specific strategy

dominant strategy profile

if every player plays a dominant strategy
can only occur in NE's

nash equilibrium NE

player can not improve payoff by unilaterally changing strategy
unilaterally means the other players don't change strategy
games can have multiple nash equilibria
if all players choose dominant strategies

mixed nash equilibrium

if fixed probability for each option is defined this exists

best response

strategy given belief about strategy of others
if best response same strategy for all options, it is dominant

price of anarchy PoA

cost(NE-) / cost(SO), with NE- as NE with smallest payoff

optimistic price of anarchy OPoA

cost(NE+) / cost(SO), with NE+ as NE with highest payoff

mechanism design

focuses on designing games where all behave nicely

auction

each bidder has secret value for good, bids for good

auctioneer sells good to bidder

truthful auction

if no bidder has incentive to bid different price than value

9.2 proofs

OPoA can be $O(n)$

assume two nodes with connection $1-\epsilon$ s

those two nodes each have n nodes with 0 cost connection

if one caches, cost is $1 + n/2(1-\epsilon) = a$

if two cache, cost is $1+1 = 2$

the OPoA = $a/2 = O(n)$ for ϵ s to 0, n to inf

Braess' paradox

assume two streets, each with fast then slow part and vice-versa

assume fast street depends on traffic (slower with more traffic)

if connection road build between the two fast parts

then total travel time slower than as if no road would exist

first price auction is not truthful

highest bidder can reduce price by $b - \epsilon > b_{\text{others}}$

bidding truthful is dominant in second price auction

do case distinction with b_{max} , b , and value v of item

under/overbidding does not change payoff, or decreases

NE for selfish caching can be done

as mechanism designer can choose payoffs

can payout if no one caches, or vice versa for all subset

tit-for-tat

always mirror the action of the other player

enforce with cryptocurrency, reputation systems

give something for free to bootstrap

9.3 games

prisoners dilemma

if both defect get high penalty

if both cooperate get low penalty

defector get no penalty if the other cooperates

if repeated, called iterated prisoners dilemma

selfish caching

each node has demand d for file

each node can cache file locally for cost 1, or ask another node

if asking other node, there is a cost d for transfer

if no one caches, cost is $+\infty$

rock-paper-scissors

best response changes with each move, therefore no pure NE

expected payoff 0 with $1/3$ p mixed NE

bidding with full payout

each bidder has to pay in full, even if he does not win

will go indefinitely because rationale is to keep bidding

either don't play, split earnings (collude), or start highest

9.4 algorithms

first price auction

sell price to highest bidder

second price auction

sell price to highest bidder for price of second highest

use concept for selfish caching

selfish caching

assume storage cost 1, transfer from node n costs d_n

choose node n with highest demand, it needs to cache

remove all nodes from set of nodes S with $d_n < 1$

repeat till S empty

9.5 exercises

pure nash equilibrium NE

write down for each entity best response strategy

find assignments which are allowed by strategies

social optimum SO

choose assignment which would be best for all

could also be a NE

price of anarchy PoA

take worst NE, then NE / SO, should be ≥ 1

optimal price of anarchy OPoA

take best NE, then NE / SO, should be ≥ 1

mixed nash equilibrium MNE

define p, q probabilities for player 1,2 to play option 1

then $(1-p)$, $(1-q)$ defines probability for player 1,2 to play option 2

write down utility for each player depending on p, q

$p(\text{payoff}_{\text{for}_q} * q + \text{payoff}_{\text{for_not}_q} * (1-q)) + \text{reverse}$

then check where payoff = 0 for both players, this is MNE

proof that it is the only one

10 Locking

10.1 general

focus

multiple faithful processes with shared memory

practical performance the most important factor

wait by blocking

let scheduler start a new thread, and pause

wait by spinning, busy-trying

keep trying to unlock by spinning on memory location

but memory changes not instant (could use memory barrier)

but memory does not guarantee sequential memory

but processors reorder instructions

10.2 testAndSet

java api

single binary value

stores true and returns value from before

getAndSet(true) is same

provokes a lot of bus traffic

TTAS

read fields repeatedly and wait till correct value read out

then use testAndSet to finalize operation

avoids some unnecessary bus traffic

but on testAndSet, invalidation storm happens at all cores

contention

if multiple threads try to acquire lock at the same time

low if few, high if many

exponential backoff

wait for random time, each try double possible range

read-write lock

single shared integer, set to -1 for write, set to >0 for read

readers can starve writers

ticketing lock

single shared integer, first part is head, second part is tail

on locking, increment tail value, set as current ticket number

on waiting, wait till head equals current ticket number

in release, increment head by one

10.3 queue locks

problems tackled

less cache-coherence traffic (no spinning on same location)

better critical section utilizing (instant control switch)

generally

enqueue thread, each spins on different location

each new thread is notified by his predecessor

first-come first-serve fairness

array-based

boolean array, entry is true if respective thread allowed

add pads between the entries to avoid false sharing of cache line

CLH queue lock

list of nodes, each node represents thread

each node has field is_waiting (boolean), pred (predecessor node)

on lock acquire, set as tail of list, set pred, set is_waiting to true

on waiting, spin on pred.is.waiting, wait for it to be false
on unlock, set this.is.waiting to false
can reuse pred node for future accesses
spinning on different core could be an issue

MCS queue lock

list of nodes, each node represents thread
each node has field is.waiting (boolean), next (successor node)
on lock acquire, set tail to self, set (old tail).next to self
on waiting, set is.waiting to true, spin till it is set false
on lock release, set next.is.waiting to false
spins on same core

queue lock with timeouts

allow timeout (max time caller wants to wait)
trivial with backoff algorithms, difficult with queues
do not remove nodes, rather mark node as abandoned
do so by setting itself as value of pred of next node
waiting node spins on in field pred
if (field is null) then keep waiting
elseif (field is static node AVAILABLE) then proceed to critical
else repeat on pred.pred ("change node")

11 concurrency

11.1 definitions

coarse-grained synchronization

acquire lock for whole data structure

fine-grained synchronization

split the object into independently synchronized components

optimistic synchronization

search with no locks, verify after talking locks

lazy synchronization

postphone work (split logical, physical removal)

non-blocking synchronization

use of atomic operations such as compareAndSet

11.2 concurrent list

11.2.1 concurrent reasoning

define invariants, show they hold after creation
show that no thread can make property false

11.2.2 freedom from interference

only methods where concurrent reasoning holds can access

11.2.3 abstract vs concrete

abstract value defines the perceived functionality
concrete representation defines the implementation

11.2.4 invariants

sentinels (head, tail) are not added/removed
keys are unique and in order

11.2.5 safety property

linearizability (identify linearization point, atomic step)

11.2.6 non-blocking properties (liveness)

wait-free, every call finishes in finite number of steps
lock-free, if some finish in finite number of steps

11.2.7 implementations

coarse-grained synchronization

every thread must acquire single lock from list
linearization point is where the lock is acquired
satisfies same progress property as lock

fine-grained synchronization

hand-over-hand locking, always pred & curr is locked
all threads must acquire in same order to guarantee progress
starvation free because of lock and no deadlock possible

optimistic synchronization

search without acquiring locks, then lock & check
needs to traverse a second time in validate
not starvation-free

lazy synchronization

add boolean to node called marked
contains() wait-free because can simply traverse list

add() locks, checks for not marked, then adds
remove() lazy, sets marked to true, then changes pointers
add()/remove() not starvation free
linearization on setting the mark

non-blocking synchronization

cannot use compareAndSet (two following nodes removal)
use other api which includes marked bit in reference
needs to take clean-up approach to avoid 2nd traverse of list
remove() marks elements as removed, then tries once to remove
add()/remove() lock-free, contains() wait-free
add()/remove() look for marked, clean up & restart if found

11.3 concurrent hashing

11.3.1 modes

open addressing (hash refers to single item)
closed addressing (hash refers to bucket of items)

11.3.2 properties

easy to parallelize (disjoint-access parallel)
deal with collisions (different items, same hash)

11.3.3 approach

define base hash class with contains, add method
abstract acquire (before contains, add)
abstract release (after contains, add)
abstract policy (checks if resize should be called, after add)
abstract resize (after policy returned true)

11.3.4 implementations

course grained hash set

acquire(), release() use simple lock
policy() checks for 4 average, resize() doubles buckets

striped hash set

each bucket has own lock, lock array size not changed
resize() acquires all locks in specified order
each lock is responsible for all buckets $b \bmod n$

refined hash set

introduce reference with owner, boolean
owner can resize locks, boolean indicates if resize in progress
acquire() needs to check after locking if the locks still valid
resize() sets itself as owner, acquires all locks, then resizes

12 Consistency & Transactional memory

12.1 concurrent objects

sequential object method specification

precondition (state which has to hold)
postcondition (return values, exceptions)
side effects (what happens when calling)
must assume meaningful state only between calls of method

keywords

program order (method execution order of single thread)
compositional (system fulfils if all objects fulfil)
execution in isolation (no other threads take steps)

progress conditions

non-blocking (any pending invocation has correct response)
blocking (one thread can deny others progress)
wait-free (invocation finishes in finite number of steps)
bounded wait-free (like wait-free, but bounded steps)
population-oblivious (performance not dependent on thread count)
lock-free (some invocations are wait-free)

dependent progress conditions

progress occurs if lower layer provides some guarantees
lock-based algorithms can only ever guarantee dependent progress
deadlock-free (some process will eventually make progress)
starvation-free (any process will eventually make progress)
obstruction-free (any isolated process finishes in finite steps)

method calls should appear

- (1) sequentially one-at-a-time
- (2) effective in real-time order if separated by quiescence
- (3) effective in program order
- (4) effective instantaneously at some moment (linearization point)

linearization point

where method effect is visible to all other users

if locks, the critical section
if no locks, reference / boolean writes or similar

quiescent

if no pending method calls ("break")

quiescent consistency

needs (1), (2)
can reorder any calls in groups separated by break
w_a(0) break; w_a(1), r_a(0), r_b(1)
is non-blocking, compositional
use in printer

sequential consistency

needs (1), (3)
can shift around calls but no reordering
w_a(1), w_b(2) break; r_a(1), r_b(2)
there may be one than one valid orders possible
processor provides memory barriers to deny reordering
is non-blocking, not compositional
use for banking

linearizability

is composable
needs (4)
is non-blocking, compositional
use for stock-trading

12.2 transactional memory

why locking is bad

priority inversion (low thread holds lock for high)
convoing (thread holding lock is descheduled)
deadlock (threads lock same resources in different orders)
maintenance relays on conventions, comments

why compareAndSet is bad

operate only on single word, algorithms get complex
can't implement multiple-compareAndSet

why compositionality is bad

can not compose multiple atomic calls to single atomic
modularity of algorithms hard

transaction

if (no error during execution) then commit else aborts
are serializable (coarse-grained linearizability)
no dead-lock, no live-lock
executed speculatively, makes tentative changes
can be nested (modularity), parent may not abort if child does
globally "locks" with all other active atomic sections

12.3 software transactional memory (STM)

12.3.1 transactional thread

has onCommit, onAbort, onValidate callables
create thread for each transaction which runs callable
retry transaction if no error till status is committed

12.3.2 transaction

has status active, aborted, committed
exposes this status to the contention manager

12.3.3 atomic objects

way for communication between transactions
needs get/set for all fields, atomic arrays
needs copyTo method for the different transactions
openRead / openWrite for getter/setter, afterwards call validate()
validate() returns false if current transaction has been aborted

12.3.4 contention manager

resolves conflicts between two transactions
can abort transaction by getAndSet to status
backoff (back-off doubling each time, if limit reached abort)
priority (priority by timestamp, abort younger, else wait)
greedy (priority by timestamp, abort younger & waiting)
karma (priority by work done, abort less work)

12.3.5 obstruction-free atomic object

setup

each object has tree fields (owner, old version, new version)
owner is last transaction which accessed the object
old version is object state before transaction arrived
new version has changes from transaction if any

determining current version

if (owner committed) then new version is current
if (owner aborted) then old version is current
if (owner active) then no current version

field access

if (owner committed) then set old = new, sets new
if (owner aborted) set new = old, sets new
if (owner active) asks contention manager

12.3.6 lock-based atomic object

concept

remove need for double reference traversal
instead lock on writing (reading doesn't need to)
global version clock, increased each time a transaction commits

object fields

stamp (clock of last time written to)
version (instance of sequential object)
lock (is a lock)

transaction

uses local read/write set of objects (with versions obtained)
does all changes to this set

validate of transaction

locks all write objects, compareAndSet to increment global clock
check all read objects are not locked, validates local read versions
updates all stamps, versions of write set objects

12.4 hardware transactional memory

cache coherence

detects synchronization conflicts
listens on the bus (snooping) to detect changes

MESI

modified (line is modified, needs to be written back)
exclusive (line not modified, and not on other cores)
shared (line not modified, but on other cores)
invalid (line invalid)

MESI load actions

if (exclusive) then write back modified, invalidate shared
if (shared) then change exclusive to shared

transaction extension

add transactional bit to cache line tag
if (bit set && modified) then do not write back
if (bit set && invalidate) then abort transaction
clear bit on committing transaction

limitations

aborts can be hardware (don't retry) or conflicts (retry)
size of transaction is limited to cache size
length limited to scheduling by OS / next cache clean-up
not fully associative cache may has impossible transactions
transactions can starve each others

13 blockchain / coins

13.1 piChain

unite block-chain (simple, fault tolerant, but eventual consistency)
with consensus protocols (strong consistency, but poor scalability)

target properties

fault tolerant (no byzantine, but crashes)
fast (no base overhead, strong consistency within one RTT)
quiet (no heartbeat needed, no work no messages)
scalable (no quadratic number of messages)
light (no explicit leader, simple architecture)
available & consistent (only short united phases needed for consistency)

transactions

executable commands
unique id (node which created transaction, sequence number)
broadcasted to whole network
every transaction will eventually be in a block

blocks

many or single transaction
unique id (node which created block, sequence number)
depth field (number of transactions including ancestors)
contains pointer to parent block (deepest block creator has seen)

node states

node is either quick, medium, slow

when quick, creates block as soon as new transaction found

when medium, wait with creating block longer than quick (wait 2t)

when slow, only create block if fast nodes somehow don't (wait 4t)

node promotes when new block created

nodes set to slow when new block seen && (deepest || creator is quick)

strong consistency

nodes can only commit block if parents are transitively committed

to commit a block, node must convince majority twice with paxos

comments on approach

can commit few times to shorten chain, or often for fast consistency

paxos adaptation

node communicates with other nodes, local storage contains only blocks

T_store called b_supp, T called b_prop, b_max called T_max

b_new is the new block, b_com is c

can skip phase 1 by directly sending propose(b_new, empty)

healthy system

if only one quick node, then paxos can run in parallel

if all slow nodes, few may become medium, but only one quick

13.2 ethereum

turing-complete scripting language to serve as base for programs

value-awareness (can withdraw fine tuned amount)

stateful (scripts can have state)

blockchain-awareness (scripts know where they are in chain)

state (accounts)

20-byte address, is target & source of transactions

nonce (counter to make sure transaction only processed once)

ether balance, contract code (if any), storage (default is empty)

can be externally owned (private keys) and contract owned

all forms of addresses have same possibilities

messages

like transactions, but with storage

if (target is contract) then can have return value

transaction

singed by externally owned account

contains recipient, signature, ether & date to be sent

STARTGAS (max amount of steps) & GASPRICE (fee per step)

execution terminates if STARTGAS empty, else rest is transferred back

state transition

validate transaction (signature, nonce, balance of sender)

subtract GASPRICE * STARTGAS, increment nonce from sender account

init GAS = STARTGAS, subtract some GAS to pay for transaction

locate receiver, transfer ether, if contract, run its code

if transfer failed because GAS = 0, then revert state changes

add fees to miner account, refund rest to sender

code execution

stack-based ethereum virtual machine code

can be written in Serpent (high level language)

32byte stack, memory byte array, key-value storage of contract

blockchain

each block stores transaction list and most recent state

check transactions, previous block, proof of work are valid

can efficiently store the state with pointers

applications

financial (contracts, wallets, payment)

token systems (property, coupons, coins)

possible implementations

namecoin (DNS like, persist name to storage of contract)

file storage (continuously pay out anyone holding a chosen file block)

organisation (execute sub-contracts as agreed on)

saving wallet (bank which can pay out 1% as security)

crop insurance (connect with weather feed, pay out when bad)

data feed (pay out all which are within 25-75% of median)

escrow (signature with different weights)

cloud computing (randomly pay out if participant is correct)

gambling (difference on block hash)

14 Exercises

14.1 asynchronous riddle

group of people each continuously enter room with switch

find out when everybody visited the room at least once

switch position known

all turn switch on exactly once

selected leader turns switch off, counts to n-1

switch position unknown

all turn switch on twice

selected leader turns switch off, counts to 2n-2