

compiler design

32362 characters in 5036 words on 833 lines

Florian Moser

August 14, 2018

1 introduction

1.1 definitions

programming language

a precisely defined language

compiler

program written in a host language

translates a programming language to another one

program

sequence of expressions executed by a stateful processor

expression

reads / modifies state & determines next expression

special E.stop halts the program

execution

elaboration of expressions on some machine M

M defined by software (virtual)

M realized by hardware (physical)

translated (further compiled down) or interpreted (directly executed)

1.2 formal specifications

operational semantics

abstract machine A interprets sequences of steps

effects produced by A determines meaning

denotational semantics

mathematical construct defining outcome

steps to reach desired outcome can be chosen

axiomatic semantics

rules that describe effects of operations (assertions on programs state)

1.3 compiler

1.3.1 reasons for translation

L1 faster to develop, no hardware exists

L2 more efficient/stable, actual hardware exists

1.3.2 tasks

preservation of semantics

same results for all legal inputs

illegal inputs, non functional requirements may differ

resource management

limited registers & storage, finite representation

respect execution environment (caching mechanisms)

constraint checking

ensure program can indeed be translated

inject code to preserve constraints at runtime

1.3.3 models

ahead of execution time (AOT)

compilation prior to execution

commonly used for languages without environments

just in time (JIT)

continuous compilation at runtime (optimize frequently used methods)

commonly used for languages with VM (java, c#)

1.3.4 architecture

frontend

read input, transform, check constraints

intermediate representation (IR)

compiler internal format (tree form or more complex structures)

expresses all language constructs/concepts

may consists of multiple stages & optimizer

backend

generate code respecting machine constraints

1.3.5 design

correctness

design for correctness, tradeoff with performance

if (correct program & valid inputs) behaves same as L1

else behaves in a defined way (by L1 / L2 / machine)

2 compiler built in course

specifics for the JavaLi compiler build in course

2.1 ANTRL

ambiguity resolution by taking document order

direct left-recursion (in same rule) allowed

2.2 IR (tree based)

tree determines evaluation order

2.3 code generator

child code must be generated first

registers

list of possible registers given by architecture

minimize usage by evaluating first left or first right

handle too few registers (abort, restructure, spilling)

template

defines behaviour of IR node or collection of nodes

contains code & describes bookkeeping tasks

leaves result in register & reflects machine properties

process

traverse tree, retrieving most specific template possible

emits code (by writing to buffer / file)

does bookkeeping as specified by template

const

put value in register

if too big then require frontend to throw error

else deal with it in assembly (add to .data, assemble at runtime)

var

if left delivers an address for parent

if right delivers a value

operators

map each operator to an assembly instruction

put result of operation in one, allow reuse of other

assign

move value from right to location at left

improve code generator

special instructions with const childs (incr for +1, shift for *4)

directly use memory operand instead of movl in register first

3 frontend

3.1 grammar

provide set of rules to generate strings

elements

terminals, characters (a,...,z, empty)

non-terminals, syntactic variables (A,...,Z)

start symbol (S)

set of productions ($Z \rightarrow aZb$)

L(G)

production \Rightarrow as left-hand side (LHS) \rightarrow right-hand side (RHS)

derivation \Rightarrow^* as repeated applied productions

L(G) is set of strings w, for all w there exists S \Rightarrow^* w

grammar type overview

context-free (type-2) if only single non-terminal on LHS

linear if context-free and at most one NT in RHS

left-linear if linear and NT at left end of RHS

right-linear if linear and NT at right end of RHS

regular (type-3) if left- or right-linear

type-2 grammar

efficient analysis techniques known

non-terminal evaluation order does not matter

can be recognised by stack machine

type-3 grammar

generate regular languages, regex

recognised by finite deterministic automata

derivations

start at S, then find next rule to get to word w ("top-down")

left-most (always pick left-most non-terminal)

right-most (always pick right-most non-terminal)

ambiguous grammar

if more than one left-most derivation can be found

avoid ambiguity by changing language/grammar or defining precedence

3.2 lexical analysis (scanner)

transform character stream into tokens for parser

token assembly

stop if character not in token anymore (like whitespace)

maximal munch strategy

put into current token as much as possible

but error reporting difficult, breaks down with complex languages

couple with parser

ask parser for expected token types (top-down vs bottom-up)

3.3 syntactical analysis (parser)

prove that word is in grammar by providing steps of derivations

can choose way to parse if multiple non-terminals

employ regular expressions

define regex for each grammar token ("top-down" approach)

because exact steps to generate symbols do not matter

can be recognised by both DFA and NFA (finite state machines)

construct parse tree

as the result of the derivation, removing unnecessary information

root is S, nodes from productions, leaves from terminals

produce IR

convert the parsed syntax tree to IR

removes intermediate steps & compresses info

like $E(\rightarrow Id) - OP(\rightarrow +) \rightarrow E(\rightarrow Id)$ transformed to $+(\rightarrow Var, \rightarrow Var)$

3.4 top-down parser

start with S, work towards retrieving w

3.4.1 simple stack machine

decision based on current input & top of stack

error/accept as output actions

reduction to change symbol on stack

match to consume from input & pop from stack

bruteforce possible as words & grammars finite

3.4.2 strategies

backtracking

deterministic bruteforce variant

start with S on stack

if (top of stack & input match) then consume input

else if (can apply production) then apply production

else if (can revert & try new production) then do it

else if (can restore input & try new production) then do it

else reject

predictive

if grammar in LL(k), the next k symbols determine action

try to construct parsing table M with (x=input, y=non-terminals)

define next production for all valid top-of-stack & k next input

combinations

entries which stay empty throw a parsing error, x=\$ and y=\$accepts

if such M can be constructed then grammar not ambiguous

3.4.3 parsing table M for LL(1)

LL-regular grammar LL(1) (left-to-right scan, left-most derivation)

compute FIRST(p)

from left to right evaluate; let current be first element of RHS

(1) if current is terminal or empty, then add to FIRST(p) and stop

(2) else add FIRST(current) - e to FIRST(p)

(2.1) if $e \in \text{FIRST}(\text{current})$ then continue

(2.1.1) if no more following elements then add epsilon to FIRST(p)

(2.1.2) else continue at (1) with the following element

repeat for all RHS variants, start with RHS fulfilling (1)

compute FOLLOW(p)

find all occurrences of LHS in others RHS

let p' be the rule with the occurrence, let match be LHS in p'

(1) if S include \$ in FOLLOW(p)

(2) if match at the end add FOLLOW(p') to FOLLOW

(3) else add FIRST(next) - e to FOLLOW

(3.1) if $e \in \text{FIRST}(\text{next})$ then continue at (2) with the following element

repeat for all LHS

construct M

write terminals to x axis, non-terminals to y axis

for each production $A \rightarrow B$ with $B \neq \epsilon$ write it to the FIRST(B) cells

for each $B = \epsilon$, write $A \rightarrow \epsilon$ to all FOLLOW(A) cells

add extra row/column with \$, writing ACCEPT in the crossing

process stack

start with stack = \$S, pointer to first entry, tree = S

(1) lookup table with M[entry, S], replace with S on stack & write in tree

(2) if top of stack = entry then pop, advance pointer, goto (2)

(3) else goto (1)

3.5 bottom-up parser

start at w, tries to get back to S

3.5.1 simple stack machine

decision based on current input & top of stack

error/accept as output actions

push to shift input on stack

reduce to replace top of stack with production

3.5.2 strategies

predictive

stack as scratchpad; with symbols to be used later

delay recording production until its obvious

need to know derivations of RHS to know if its production applies

add a control-symbol between each symbol to preserve context

use parsing table M[control symbol, t] to get next action

3.5.3 parsing table M for SLR(1)

LR-regular grammar LR(1) (left-to-right scan, right-most derivation)

create LR(0) sets with FSA automaton

build simplified LR(1) parser (SLR(1)) parser

construct FSA

dot . = processing position, augment grammar with $S' \rightarrow .S$

(1) create first state of FSA as $S' \rightarrow .S$

(2) compute closure by following NT after dot ($S \rightarrow .A, A \rightarrow .aA \mid .b$)

(3) create all possible transitions (a produces $A \rightarrow a.A$), goto (2)

give all n possible states a name I_k

create table

compute first & follow, number states & productions

states to x axis

actions (terminals+\$) & gotos (non-terminals-S') to y axis

for terminal transitions, write Shift:#state to corresponding cell

for non-terminal transitions, write #state to corresponding cell

for states with $S' \rightarrow S$. write accept to \$column

for states with $A \rightarrow b$., write Reduce:#production to FOLLOW(A) column

process stack

start with stack = <0> & set pointer to first entry

(1) tablelookup with state (right-most control symbol) & entry

(2) if shift then write entry + <shift_state> then goto (1)

(3) if reduce then record reduction $X \rightarrow Y$ & remove all from stack

(3 cont.) tablelookup with state & X, then write $X + \langle \text{goto_state} \rangle$

3.6 parser varia

generators can create LR(0,1) items, LL(1..) parsers

canonical LR (CLR)

uses LR(1) items (lookahead of one)

may resolves ambiguities, but FSA needs more states

lookahead LR (LARL)

same number of states as SLR(1) but equally powerful than CRL

limitations of context-free grammars

want to find errors with parser

$L1 = \{ a c a \mid a \in \{a, b\}^* \}$ for c separate body/definition

L1 could find undefined variables, but not context-free

$L2 = \{ a^n b^n \}$ for a definition and b calling site

L2 could check if method invocations are proper, but not context-free

use semantic analysis for that

4 semantic analysis

check properties (types & constraints)

perform transformations (casts, default values, initializers, ...)

4.1 attribute grammar

context-free grammar extended with context-sensitive info

4.1.1 attributes

attached to non-terminals

hold value

synthesised (from children) or inherited (from parent/siblings)

notation

name non-terminals of production, possibly with index ($\langle P \rangle$, $\langle P.0 \rangle$)

can add properties functioning like global variables ($\langle A \rangle.N$)

can specify calculations ($\langle A.0 \rangle.N = \langle A.1 \rangle.N + 1$)

can specify conditions (iff $\langle A.0 \rangle.N == \langle A.1 \rangle.N$)

4.2 symbol table

repository of program symbols

nesting may hides symbols (variable hides field)

enables checks for classes, fields, methods, variables

4.2.1 structure

global symbol table

contains class entries

contains constants such as true/false, void, ...

class entry

fields (size, location (offset from this))

methods (location)

base class

method entry

parameters (order, type)

variables (size, location (offset from frame pointer))

debugging

definition location

4.2.2 construction

walk over AST, create entries

field/method references

implicit "this" reference if not specified

explicitly specified class reference

perform checks

no duplicates (class / methods)

no undefined fields/method/classes (resolve reference)

valid inheritance (no cycles, default extend Object)

valid type (parameter passing, variable assigning)

valid type conversions (add int + float = float)

valid array access ($0 \leq i < A.length$; might needs intermediates)

defined program start (main method)

no unused variables / fields

functions return value on all execution paths

protection system respected (private/protected access)

interface/abstract correctly implemented

delayed checks

checks which can't be performed by simple analyzer
model execution checks, runtime checks, undecidable

4.3 delayed checks

model execution checks

variables initialized before usage

exceptions are caught

return statement is reachable

end of program is reachable

methods are called at least once

time/power policies respected

undecidable checks

loops/program terminates

runtime checks

array out of bounds access

implicit null checks by pointing to low & protected memory region

triggers protection error (SIGSEGV); VM catches & throws exception

5 software engineering

5.1 patterns

strategy

encapsulate set of related algorithms into class hierarchy

interface (IStrategy) passed to context (IConsumer)

ConcreteStrategy implements IStrategy

specialization

create new child class with more specific functionality

child inherits parent class, overriding or adding methods/fields

visitor

performs a double invocation (caller/callee determined at runtime)

ouples data structure w/ unrelated functionality (IElements)

with access methods w/ unrelated operations (IVisitor)

IVisitor with visitMyElement implemented by MyVisitor

IElement with accept(IVisitor) implemented by MyElement

MyElement implements accept() { visitor.accept(this) }

MyVisitor has method for each IElement implementation

the visitor or the element traverses (possibly tree) structure

observer

presentation (IObserver) is notified about change in data (ISubject)

ISubject implements attach(Observer), detach(Observer), notify()

IObserver implements update()

subject calls update() upon change; then observer gets state of subject

collection extensions

provide hook for externally defined functions for easy extensibility

high order function f (named such because passed as an argument)

MyCollection provides hook(IFunction f) which applies f to all elements

MyFunction implements IFunction with map(IElement)

5.2 software quality

general approach

identify quality dimensions

set priorities, discuss tradeoffs

establish objectives (time, cost, resources, performance)

perform measurements to fulfil objectives

quality assurance

follow guidelines of software construction

pursue appropriate testing strategy

informal reviews, formal reviews, external audits

internal characteristics

primarily interests developer

flexibility (can be modified for unintended environments)

reusability (can be used in another execution environment)

testability (test if requirements are fulfilled)

maintainability (address changing needs easily)

readability (consistency, conventions, structure, used constructs)

understandability (separation of concerns)

external characteristics

primarily interests user

correctness of specification, design, implementation

robustness to handle invalid input / environment

usability (fit for purpose)

adaptability (incorporates with other systems)

integrity (access control)

traceability (access traces / logs provided)
accuracy (results as expected)
reliability (performs under assumed environment)

5.3 informal reviews

execute tests & understand source code to catch errors early
explicitly not to assign blame or evaluate staff/project

5.3.1 walk-through

devs present system

5.3.2 code reading

feedback via comments, done by co-dev

5.3.3 inspection

reviewers look at code, then meet with dev

roles

moderator (by technical person, guides discussion)
author (by author, explains decisions & gives overview)
reviewers (point out good/bad & possible problems/defects)
scribe (possibly by moderator, records discussion)
management is absent

process

planning (code sent to moderator, selects reviewers)
setup (short overview given by author)
preparation (code & documentation is read)
meeting (factual questions are discussed)
report (scribe publishes report)
follow-up (moderator assigns problems to repair)

5.4 openJDK example

review process

develop & test code
upload patch to server
inform mailing list of patch
reviewers inspect code & give feedback
push change set (code + description + reviewers)

roles

higher accepted changesets (#) unlocks new capabilities
(0) contributor (has signed contribution statement)
(2) author (has own username & listed on project webpage)
(8) committer (push-right & sponsors others change sets)
(32) reviewer (decides if code is good enough)

6 code generation

primary requirements are correctness & performance

6.1 key activities

code selection (decide assembly for IR nodes)
code scheduling (determine order of execution)
register allocation (decide which operand in register)
register assignment (decide which register used)

6.2 x86 rules

one operand may reside in memory or intermediate
6 int registers available

calling conventions

caller preserved %eax, %ecx, %edx
callee preserved %ebx, %ebp, %esp, %esi, %edi
callee must restore %ebp %esp

6.3 assembly

assume a in %eax, b in %ecx

```
movl $2, %eax # save 2 to b
movl %eax, %ecx # save a to b
addl %eax, %ecx # b = a + b
incr %eax, # a++
notl %eax, # a = - a - 1
xchg %eax, %ecx # temp = a; a = b; b = temp
imull %ecx, # a = b*a (result in edx:eax, each 32 bits)
imull %eax, %ecx # b = a*b (32 bit result)
cmpl $2, %eax # set 0 flag if equal
```

```
je LABEL_1 # jump if cmp before equals
jg, jge LABEL_1 # jump if first greater (+equal)
jl, jle LABEL_1 # jump if first lesser (+equal)
jmp *(%eax) # indirect jump
jmp (%eax) # jump to address
LABEL_1:
call LABEL_1 # push address of next instruction to stack, then jump
```

6.4 sample assembly

```
push %ebp # preserve %ebp
mov %esp, %ebp # top of the stack is now in %ebp
sub $16, %esp # make stack frame bigger
push $0xDEADBEEF # push stuff on stack; increasing %esp+4
mov 4(%ebp), %eax # get 1st passed parameter %ebp
mov %ebp, %esp # restore original %esp
pop %ebp # restore original %ebp
ret # returns (equivalent to pop %eax jmp *%eax)
```

6.5 operand access

constant

put in register or use directly
enforce size contains

fields

reference of object + offset from symbol table
use resulting memory location to read/write

arrays

reference of array + blocksize * index
row-after-row called row-major (x[0][0], x[0][1], ...)

access

evaluate expressions & in intermediate
evaluate from left to right

6.6 assignment statement

assign to variable, field, array
handle register shortage; preserved on stack & override

6.7 conditional statement

if-then

eval/test condition
if true then continue, else jump to end

if-then-else

eval/test condition
continue if true then jump to end, else jump to else

6.8 loop statement

while expr

eval/test condition
if true then continue then jump back, else jump to end
but one unconditional branch for each expression

while expr v2

jump body
eval/test condition
if true jump back to body

6.9 method invocation

call site from caller invokes callee on target with actual arguments

calling conventions

decides what caller/callee prepares
decides where date, return address, return value are saved

call actions

identify target & starting address of callee
handle parameters (evaluate, push into stack)
save caller saved registers
find space of temporaries
determine & store return address
find space for return value
set up activation record for callee
transfer control to callee

return actions

restore registers
transfer control to caller
deliver return value
remove activation record of callee

6.10 heap

eats up, from low to high addresses
contains object instances

6.11 stack

eats down, from high to low addresses
setup by run-time system, some assembly instructions

structure

frame/base pointer fp determines start of frame (high address, low)
stack pointer sp determines end of frame (low address, high)
return value, parameters + target, return address (caller)
old fp, old sp, locals, temps (callee)

access stack elements

offsets persisted in symbol table
access parameters with fp + offset (go to caller area)
access locals with fp - offset (inside callee area)
access temps with fp - offset (can grow indefinitely)

6.12 object layout

object referenced by pointer p
p + 0 pointer to v-table
p + 4,... the corresponding fields

v-table

created completely for each class
method order in same order as base class
code location same as base class if not overridden

call method

dereference object pointer (movl (%eax), %eax)
dereference v-table at this + 0 (movl (%eax), %eax)
call code location from v-table offset (call *4(%eax))

7 data flow analysis

correct & accurate for all possible inputs (conservative)

7.1 applications

reason about correctness (ensure value assigned)
enable optimizations (compile-time expression evaluation)
useful for tools (find all references of symbol)
remove unused code

7.2 control flow graph (CFG)

constructed over single method
basic box is executed as a unit

construction

create blocks (started by label, ended by return/goto/condition)
reformulate if statements to "TCond# = condition; if (TCond#)"
connect blocks which are executed after one another
name blocks in order like B#
add ENTRY block (pointing to first block)
add EXIT block (successor of final blocks like returns)

analysis

local, inside basic block (like subexpression elimination)
global, over CFG (like null checks)
inter-procedural (across methods/functions)

point

after specific statement s1, before s2
after specific basic block b1, before b2
make claims at specific points

paths

for all pairs xi, xi+1 it holds one of these two
(a) xi is point before b, xi+1 is point after b
(b) xi is point after b, xi+1 is point before connected b2
infeasible to process all (possibly infinite) paths

7.3 side effects

never

local variables / parameters
field assignments (if no dereferencing needed)

in expressions

null-dereference

non-pure method evaluations
invalid array access

7.4 data flow framework generalization

D as direction of data flow (forward/backwards)
V domain of values (possible results + bottom & top)
|><| meet operator (operating on lattices)
F transfer function (process block)

termination

stop as soon as nothing has changed anymore
show that max number of transitions per element exists

7.5 constant propagation

replace constant variables with its values

process in block

set all variables to bottom
go to specific value if const assigned
else go to top

transfer

overapproximate reverse-T \rightarrow exact \rightarrow T

evaluation

directly place constants in code

7.6 reaching definitions

check if/where variable is defined which is used in statement
d reaches u if d is not killed on the path to u

process in block

create table, for each variable all definitions
gen by taking last definition in block of variable
kill with all other definitions of variables in gen
overapproximate gen = {}, kill = {all}

transfer

IN(B) = { d reaches B } = U OUT(predecessors B)
OUT(B) = { d leaves B } = gen_b U (IN(B) - kill_b)
start with OUT(start) = empty, overapproximate with all
iterate by determining IN for all, then OUT for all, repeat

evaluation

directly place expression of definitions in block
if d2 \in IN(B) can replace d2 value at its usages

uninitialized variables

let OUT(start) = { one fake definition for each value }
if variable is used with fake definition in IN(b), then possibly undefined

7.7 live variable

check if variable is used after some assignment
dead variables need not to be stored / computed

process in block

def set with all var defined in B prior to any usage
use set with all var used in B prior to any definition
overapproximate def = {}, use = {all}

transfer

OUT(B) = { var live after B } = U IN(successors B)
IN(B) = { var live before B } = use_b U (OUT(B) - def_b)
start with IN(end) = empty, overapproximate with all
iterate by determining OUT for all; then IN for all; then repeat

evaluation

if val not in OUT(B) do not need to store/compute it

local analysis

start with OUT(B) as set of live variables
then for each statement add new live variables

use-def / def-use chains

for use of variable, persist all possible definitions
for definition of variable, persist all possible uses
do not forget loop back uses

7.8 available expressions

available if evaluated on every path and no breaking definition afterwards
breaking definition if any element is redefined
but allowed to yield different values depending on path

process in block

gen set for evaluated expression, no breaking definition
kill set for breaking definition & no following evaluation
overapproximate $gen = \{\}$, $kill = \{\text{all}\}$

transfer

$IN(B) = \{ \text{expr reaches } B \} = \text{intersect } OUT(\text{predecessors } B)$
 $OUT(B) = \{ \text{expr leaves } B \} = \text{gen}_b \cup (IN(B) - \text{kill}_b)$
start with $OUT(\text{start}) = \text{empty}$, overapproximate with all
start with $OUT(\text{other}) = \text{all}$, overapproximate with none
iterate by going top to bottom, IN then OUT; then repeat

evaluation

if expr in $IN(B)$ then do not need to recompute

8 register allocation

8.1 core tasks

which operand resides in register (map operands \rightarrow virtual registers)
which register to use for operand (map virtual register \rightarrow register)
reduce complexity by assigning operands first to virtual registers
in second step assign the virtual registers to physical ones

8.2 live ranges

range where virtual register is live (pending use)
can compute liveness of VR like any normal variable

calculation

instructions as list, for each row compute L and R
L liveness (not-killed variable v used in current/following rows)
R reaching (not-overridden definition d.v made in before rows)
live range where $v \in L$ intersects with $d.v \in R$
multiple d.v for same v possible \rightarrow multiple live ranges possible
combine liveness with reaching to decouple different definitions

granularity

starts at left of assignment, ends on right of other
or whole statements (min-length therefore 2 lines)
or whole basic block (min-length therefore 2 blocks)

interference

if ranges overlap, then interfere with each other
can model as graph; nodes are live ranges, edge iff overlap

8.3 graph coloring

k-colorable if k colors enough such that none touch
compilers try several K-colorings, then choose the best one

K coloring kempe's algorithm (simplification step)

push nodes with $\#neighbours < K$ on stack & remove from graph
until $\leq K$ nodes left (may not be possible \rightarrow graph surgery)
color the remaining nodes (trivial bc $\#neighbours < K-1$)
now pop from stack, append to graph & assign color
but not guaranteed to succeed

graph surgery

K coloring may still possible, but hard to find
if not found then heuristically "spill" (remove) node
repeat until
all $\#neighbours < K$ requirement reached
then continue kempe's algorithm

pick spill victim heuristics

at random, lowest cost estimate, lowest use count
estimating spill cost by estimating basic block usage
heuristics (loops execute 10 times), runtime info (if JIT)

8.4 spilled variables

variables residing in memory bc not enough registers
constraint that at most one operand can be in memory

keep spare register

spilled variables put into spare register upon usage
hence enforce a K-1 coloring (or K-2 if both must be in registers)

introduce temporary

introduced for each occurrence of spilled variable
filled before instruction with value from stack
lives only for two lines (load & use)
rewrite IR accordingly, rebuild graph and retry coloring

splitting

reduce spilling (temporary loads) by splitting live ranges
heuristics to know where to split

8.5 color improvements

coalescing

remove extra copy instruction if not needed
(last usage of v1 assigned to v2 \rightarrow could possibly rename v2 to v1)
coalesce by combining nodes (but maybe K coloring then impossible)
copy property edges by marking them to attempt to use same color

pre-color

depending on instruction specific registers must be used (like imull)
per-color operands; not removed during simplification

8.6 phase coupling

code selection (imull instruction)
code scheduling (determines order of execution)
register allocation (bounds $\#operands$ in registers)

9 multiple inheritance

9.1 which base symbol to use

error if no guidance provided
add parameter which defines which base class is used
require cast to either A or B
use order of extends clause
invoke both methods / randomly select one
force symbol renaming
prohibit MI usage

9.2 diamond problem (AB extends A,B; A/B extends X)

referencing X symbols ambiguous
local changes to class may break behaviour at other places

9.3 object layout problems

example

```
AB extends A,B { a, b, method() {} }  
B { b, method() {} }  
A { a, method() {} }
```

overriding fields

offset of base classes may not match to compiled functions
therefore can not share v-tables
choose to keep code of A (fields of A with same offset)
recompile B methods (bc fields have different offset)

casts with v-table duplication

offset of methods does not match, need to divide object layout
first part like A layout, referencing to AB v-table
second part start with reference to B' v-table, then the fields
when casting, offset this pointer (-0 if statically A/AB, -8 if B)

casts with pointer adjustments

add additional this offset entry in v-tables
no code recompile needed anymore
instead adapt this pointer as needed before invocation
but runtime performance suffers (at every method call)

9.4 C++

can extend with public virtual X
virtuals constructed only once in object hierarchy

10 attacks

10.1 control flow hijack

hijack control flow; continue at attacker specified location
possibly arbitrary machine code or reusing existing code
corrupt code pointers (return address, function pointer, vtable)

10.1.1 defences

non-executable data (NX)
data execution prevention (DEP)

10.1.2 return address

"ret" implicit "jmp return pointer"
try to corrupt return pointer (point to attacker location)
must reuse already existing code (because of NX/DEP)

return oriented programming (ROP)

use code snippets ending with ret, chain them together
put on stack address1, value, address2, value, ...

address space layout randomization (ASLR)

mapping program addresses to hardware
randomize hardware addresses (stack, heap, memory base)

10.1.3 stack canaries

adds data before %ebp/return address
if data is overwritten; then buffer overflow happened
but effectiveness depends on secret & runtime slower

10.1.4 powerful attacker

no data execution / code corruption (NX/DEP)
but arbitrary data modification assumed

control flow integrity (CFI)

ensure control flow stays within CFG
bc attacker needs to change CFG to execute arbitrary code
but too conservative, needs recompilation

dynamic CFI

battles drawbacks of compile time CFI
constructs CFG at run-time, supports dynamic libraries

11 Java HotSpot VM

11.1 java execution architecture

hardware, OS, JVM, Java SE/EE, Spring, user applications

why VM

non-VM languages reimplement compiler, libraries, debugger for each OS
VM implemented for all OS/architectures, more portable code

compilation

java/scala/javascript/python to bytecode (ahead of time, using javac)
bytecode executed on abstract stack machine (by VM)

execution

template based interpreter maps bytecode to machine code sequence
JIT compilers produce machine code from bytecode
both compiled/interpreted access same stack/heap
garbage collector manages heap

JIT compilers

C1 (limited optimizations, but fast compilation & small footprint)
C2 (aggressive optimizations, but high resource demands)

compile methods

because compilation expensive specific appropriate methods to compile
interpreters collect profiling infos about methods
aggressive & optimistic compilation by C1/C2 stored in code cache
if guards assumptions proven wrong then interpreter takes over again

example optimization

inline method code; guard ensures object of the expected class

tired compilation

compile interpreter (fast start), C1 (fast compile), C2 (fast code)
collect 100 samples with interpreter, then execute C1 optimization
then collect 1000 samples with C1-compiled code, then C2 optimization
but then naive code cache inefficient (high fragmentation, bad locality)

11.2 code cache

continuous chunk of memory, stores code generated by JIT compilers

types of code

non-method code (small, immortal lifetime)
C1 profiled code (medium, limited life time)
C2 non-profiled code (large, long lifetime because expensive compilation)

segment

divide code cache into non-profiled, profiled, non-methods
better cache locality (ITLB misses decreased)
sweeper (GC for code) faster because optimized for each segment

11.3 compact strings

reduce memory footprint of chars (may occupy up to 45% of data)
save as latin-1 (1 byte) instead of UTF-16 (2 bytes) (best effort)
changes GC (deduplicates strings), String (payload byte array)

11.4 ahead of time compilation

compile to native code prior to VM launch (improves startup)
compiled code shared between VMs (C1 performance)

11.5 value types

immutable, identity-less aggregates (user defined primitives)
no identity, all fields final, no subclassing

advantages

better locality (no indirection, no header)
relaxes heap pressure due to stack allocation (less GC)
new possible optimizations (like scalarization)

minimal value types (MVT)

subset of value types, meant as early access
denote class with Value Capable Class (VCC) annotation
then JVM derives value type class (DVC)
new value type bytecodes / method handles

storage formats

on heap, with header
in thread local value buffer (TLVB) with header
scalarized by JIT code without header on stack/registers)
flattened array/field with type infos stored in container header

flattening

remove references by inlining value types
for fields (only non-static)
for arrays (respect long alignment)
GC needs support to find references in inlined value types

optimizations

copy detection (reuse existing)
passing flattened value type (passing fields as arguments)
return flattened value type & class
method handles need additional handling

explorations

currently only limited set of unit tests
check performance with runtime checks (legacy code, MVP code)