

# Object Databases

51161 characters in 6965 words on 1408 lines

Florian Moser

June 13, 2018

## 1 introduction I

### approach a new project

show the problem first  
identify MVP (whole solution stack)  
then present a solution

### product owner

lives for project  
has background infos (technology stack, DevOps)  
able to define iterations (specifically MVP)  
supports other developers (chooses the right architecture)

### basic implementation tips

show how UI is supposed to look with mockups to customer  
model database close to real world, as this does not change

### OO vs database

class vs relation  
object vs row  
value vs cell  
attribute vs domain

### programming paradigms

declarative programming (functional, logical)  
imperative programming (procedural, object-oriented)

## 2 introduction II

### 2.1 motivation

simplified application development  
OO data should be fully persistable  
database should benefit OO advantages  
simplify design & evolution  
no translation at run time

### 2.2 orthogonal persistence

#### type orthogonality

all objects can be persisted in any way

#### persistence by reachability

identifying persistent objects not related to type system  
lifetime of object determined by reachability of root elements

#### persistence independence

manipulation of long/short term data looks same

### 2.3 persistence strategies

by inheritance (extend base persistent class)  
by instantiation (with construction pass persist infos)  
by reachability (if reachable from other persistent object)

### 2.4 SQL

#### data definition language (DDL)

definition of data models  
relations, attributes, keys  
modify tables, attributes, keys  
"class definition"

#### data manipulation language (DML)

creating & management of data  
relation tuple, predicates  
execute inserts, updates, deletes  
"writing of values"

#### query language (QL)

concepts supporting retrieval of data  
query, predicate, query result  
use projections, selections, joins  
"reading of properties"

### 2.5 types

#### 2.5.1 specification

properties (attributed, relationships)  
operations  
exceptions

#### interface

defines abstract behaviour

#### class

defines abstract behaviour & state

#### literal (struct, int)

defines abstract state

#### 2.5.2 implementation

language binding, multiple per specification

#### representation

data structures motivated by abstract state  
instance variable for each abstract variable

#### methods

procedure bodies motivated by abstract behaviour  
possibly private methods not included in spec

#### subtyping (relationships)

is-a (behaviour)  
extends (state & behaviour)

## 3 OMDG standard

### 3.1 object management group (OMG)

tools & architecture for OO development  
distributed object management  
create unified modelling language UML  
informal standards body for major vendors  
promotes portability/interoperabilities  
does not develop products

### 3.2 object data management group (ODMG)

complementary to OMG  
adds data management support  
object definition language (ODF)  
object query language (OQL)  
bindings to java, smalltalk, C++

#### object model

based on OMG object model  
objects have unique identifier, literals do not  
state defined by properties  
behaviour defined by operations  
categorization possible, common properties & operations

#### extent

set of all active instances of type

#### collections

set (unordered, no duplicates)  
bag (unordered, with duplicates)  
list (ordered, insert elements)  
array (ordered, replace elements)  
dictionary (key-value store)

#### collection operations

subset only for sets  
union, intersect, difference only for bags, sets

#### relationship

many-to-many (collections for source, target)  
many-to-one (collection for source, class inverse)

one-to-one (class for source, target)  
no ternary

**persistence**  
by reachability  
database exposes root objects, schema, types

**other concepts**  
database operations, locking, transactions, metadata  
built in dates, times, intervals

**object definition language (ODL)**  
supports constructing ODMG models  
compatible to OMG interface definition language (IDL)  
defines name, extent, exceptions, attributes, relationships  
define method signatures (to be implemented later)

**object query language (OQL)**  
looks like SQL  
can use path expression like me.neighbour.name  
not complete, no explicit update  
returns bag, set with distinct, list with order by  
can use subqueries  
aggregation for collections (AVG, SUM, MIN, COUNT)  
union, intersection, except for sets  
can flatten collections

**OQL example**  
select values from collections where condition  
where exists b in books: b in a.books  
where for all p in a.authors: p.year < 1000

## 4 object oriented databases (OODB)

connection between OO world & database  
avoids mismatch between objects/relations  
provides a uniform data model  
combines features & properties of both worlds  
managing object data, making OO languages persistent

### 4.1 OODB manifesto

define features, optional characteristics, open choices  
but important properties missed (according to relational db people)

### 4.2 features of OO systems

**complex objects**  
complex objects formed from simpler ones, constructors  
tuples (represent entities & attributes)  
sets (collection of entities)  
lists (capture order)  
need transitive retrieval, CRUD, copy  
need identity (same GUID), equality (same state)  
db only knows sets, atomic tuples

**identity**  
if GUID or state matches then equal  
sharing with references  
shallow & deep equality

**encapsulation**  
object data & methods implement interface  
state only modified through public interface  
data exposed for declarative queries

**types**  
define object properties, static safety  
object structure & behaviour, separated from interface

**generalizations**  
enable better modelling, reuse, semantic complexity  
inheritance (attributes & methods from superclass)  
specialize or generalize

**inheritance**  
substitution (more operations, based on behaviour)  
inclusion (based on structure)  
constraint (like inclusion, sub is constraint on super)  
specialization (sub contains more specific info)

**method overriding**  
redefine, specialize method in sub

**method overloading**  
multiple versions of method exist

**late binding**  
appropriate method call is selected at runtime

**computational completeness**  
any computable function expressible

**extensibility**  
devs can add native types to database

### 4.3 features of db systems

**persistence**  
data survives program execution  
orthogonal implicit persistence  
orthogonality of type system & persistence (all can be saved)

**efficiency**  
secondary storage (indexes, buffers) optimize

**concurrency**  
multiuser, serialization  
atomicity, consistency, isolation, durability

**reliability**  
resilient to user, software, hardware failures  
transactions, snapshots, logging

**declarative query language**  
high level (conciseness of non-trivial tasks)  
efficient execution (can be optimized further)  
application independent (any database may execute it)  
say "what" not "how"

### 4.4 optional characteristics

multiple inheritance, type checking/inference  
distribution, versions, long/nested transactions

### 4.5 open choices

program paradigm (declarative vs imperative)  
representation system (sets, lists, more?)  
type system (generics?)  
uniformity (type, method as objects?)

### 4.6 more tools

database administration  
view definition / view derived data  
objects with roles (dynamically add/remove, like traits)  
database evolution (migrate gracefully)  
set integrity, semantic, evolution constrains  
define, modify & enforce constrains

### 4.7 technical overview

**architecture**  
vary greatly, not simple server-client model  
all do caching, locking for query/transaction processing  
need lifecycle management for objects

**granularity**  
per object/page/container  
different runtime / performance characteristics

**querying**  
try to remove join concept and replace with relationships  
performance depends on execution place, flexibility of query, indexes

**identity management**  
for relationships, uniqueness  
physical for fast dereferencing but limited flexibility  
logical for immutability but needed access translation

## 5 object relational mapping

### 5.1 mappings

db ↔ program (interface)  
db ↔ world (enterprise modelling)  
program ↔ world (simulation)

### 5.2 problems

object identity (serial/read out produces duplicates)  
multiple models (db, program, world)  
impedence mismatch (map OO to database)

implement transformation (design time)  
execute transformation (run time)  
application-specific transformation (therefore reimplement often)

### 5.3 relational database

ok for small applications  
not ok for large with inheritance, multi-value attributes

### 5.4 object relational mapping (ORM)

map OO to relational model  
persistence-related tasks already implemented  
persistence API (java, annotations define mapping)

#### set up ORM

top-down (oop  $\rightarrow$  mapping  $\rightarrow$  rdbms)  
bottom-up (oop  $\leftarrow$  mapping  $\leftarrow$  rdbms)  
inside-out (oop  $\leftarrow$  mapping  $\rightarrow$  rdbms)  
outside-in (oop  $\rightarrow$  mapping  $\leftarrow$  rdbms)

### 5.5 hibernate

java to sql, lots of backends supported

#### requirements

ensure there is a non-argument constructor  
create mapping (id, property, set (key, relationship))

#### configuration

connect to database, specify name & DBMS

#### session factory

factory = new Configuration().configure().buildSessionFactory()  
factory.close()

#### session

session = factory.openSession()  
session.beginTransaction()  
save, close, query, update, delete  
session.beginTransaction()  
session.createQuery("FROM persons")  
session.getTransaction().commit()  
session.close()

#### associations

unidirectional, bidirectional, ordered  
one-to-one, many-to-one, many-to-many

#### inheritance

table per class hierarchy (sparse, subclass)  
table per subclass (duplicated fields, joined-subclass)  
table per concrete class (no abstract, union-subclass)  
strategy can be defined for each part of application

#### annotations

replace xml with annotations, same keywords

#### solved problems

object identity  
implementation transformation  
application specific transformation

#### remaining problems

impedence mismatch  
multiple models  
transformation at run time  
annotation maintenance (new)

#### discussion

good at design time, bad at run time

## 6 android

#### application model

activities for UI  
services for computation  
providers for data  
intents request use of application components  
manifest files exposes components & defines start activity

#### lifecycle

activity reacts state changes

#### data management

SQLiteOpenHelper creates SQLiteDatabase  
Cursor c = db.rawQuery("SELECT \* FROM persons")

c.moveToNext(), c.isLast(), c.getString(2), c.getInt(3)  
ContentValues object for CRUD (simple map)

#### content resolver

reacts to URIs like content://ch.acm.personprovider/pictures/12  
content resolver invokes correct provider

#### content provider

encapsulates data management, can be exposed  
CRUD, getType() methods, uses Cursor, URI, ContentValues

#### example location

multiple providers (GPS, WLAN)  
listener waits for updates  
manager chooses best provider, has last known location

#### lack of

orthogonal persistence (type orthogonality, independence, identity)  
completeness (definition not stored; versioning problems)  
scalability (entire object graphs need to be persisted)

## 7 db4o

object-based architecture w/ physical identity

### 7.1 meta

open source native object database for .NET, java

#### key features

no conversion/mapping needed  
no changes to objects to be persistent  
local or client/server mode, single line of code  
ACID transactions  
caching & integration in native garbage collection

### 7.2 architecture

file/in-memory database  
I/O adapter  
ACID transactional slots

#### object part

marshaller  
reference system, reflector, class metadata  
api

#### query part

class/field index b-trees  
index query processor  
SODA query processor  
native & SODA queries, query by example

### 7.3 object container

connection to database  
unit of work; owns one transaction  
manages object identities, loads/unloads objects  
starts new transaction after commit/abort  
commits implicitly after closing container

#### persist

store(), delete()  
on create persistence by reachability  
on update depth is 1 per default, only primitive values  
on delete no cascade per default  
objects linked by weak reference  
config.common().objectClass(Author.class).cascadeOnDelete(true)

#### retrieve

by example (findBy, pass partly filled out POCO)  
native (using predicates)  
soda (decend("year"), constrain(class), uses the graph)

#### consistency

refresh() syncs DB state, call after delete()  
database-aware collections

### 7.4 transparent persistence

let objects implement Activatable interface  
objects bound to framework on retrieve  
modify at will, then call commit to save changes

## 7.5 activation

only loaded to certain depth, fields set to default otherwise  
occurs on collection element access, or explicit activate()  
depth tradeoff is manual activate() vs heavy memory usage

### transparent activation

on property set, activate property  
database registers itself on instance creation  
byte code insertions does this automatically

## 7.6 transaction isolation levels

read uncommitted (even uncommitted values can be read)  
read committed (only committed values can be read)  
repeatable read (always same value read)  
serializable (possible serial execution order)

## 7.7 transactions

thread-safe, but single-thread core  
no data loss, automatic recovery on system failure  
rollback discards changes, call refresh() to clean up  
read-committed (only committed values can be read)

### collision detection

peekPersisted() to get unbound instance of db version  
get read committed or stored values (configurable)

### collision avoidance

db.ext().setSemaphore(GUID, MAX\_TIME\_WAIT)

## 7.8 discussion

no impedance mismatch

### orthogonal persistence

independence (yes)  
data type orthogonality (yes)  
identification (yes)  
but explicit store/retrieve application logic

### issues

depths of activation/deletion/update is a new burden  
lack of synchronization on delete/update  
"transparency" contradicts type orthogonality

## 7.9 configuration & tuning

### defragment

remove unused fields, classes, meta data  
compact database file

### statistics

log query behaviour & performance  
log IO/network activity

### log

get all objects & classes in db file

### indexes

optimize query evaluation  
tradeoff between query & modifications  
B-trees on single fields; automatic or explicit

## 7.10 speed tuning

### object loading

configure activation depths (less loaded initially)  
use multiple containers (less complex containers)  
disable weak references (no lazy loading)

### database tests

disable database schema change detection  
disable testing of classes at startup (less validation)

### query evaluation

set indexes appropriately  
optimize native queries

## 7.11 distribution

### embedded mode

clients use same VM with db file  
direct file access (single thread, same user)  
client session (multiple threads, same user)

### client/server mode

clients use multiple VM, connect to server with db  
can only use methods from object containers  
use "out-of-band" signalling to transfer other messages

### replication

multiple servers, redundant copies  
snapshot (periodically, single-master)  
transactional (operation based, immediately)  
merge (clients send changes to master, periodically or instant)  
developer defines replication with masters/slaves

## 7.12 replication

separated from core  
bridge between db4o & relational databases  
uni/bidirectional replication between hibernate, db4o  
transfers data between providers

### steps

configure to use UUIDs & commit timestamps  
create replication object & define conflict handler  
configure direction (bidirectional is default)  
call replicate(myObj), commit() on each object (transactional)  
call objectsChanged and then replicate(iterator) (snapshot)  
call close()

## 7.13 callbacks

event triggers (activate, deactivate, new, update, delete)  
can prefix called before, on prefix called afterwards  
implement as much as needed

### use case

record/prevent updates  
check integrity / set default values  
create indexes dynamically if used often

## 7.14 control object instantiation

can be configured per class/project

### bypassing constructor (default method)

if no constructor works, if framework supports

### constructor usage

all constructors are tested with default values  
first one which works is used from now on  
if none found, object is not persisted

### translator

implement ObjectConstructor interface  
convert object to custom entity (like Object)

### type handlers

like translators, but at lower level  
translator with byte arrays, handler converts objects

## 8 Versant

object-based architecture w/ logical identity

### 8.1 meta

commercial OODBMS  
object database for java

### company

now owned by actian  
market leader in ODBMS  
telecom, military, financial, transportation

### architecture

RAID, NAS, raw devices/filesystems  
virtual system layer  
versant server (logical, physical log file)  
versant network layer  
versant manager  
C, C++, java interface

### languages

java versant interface (JVI)  
versant query language (VQL)

## 8.2 dual cache

client has object cache  
server has page cache (for each db)

## 8.3 database volumes

system volume (class descriptions, object instances)  
data volumes (increase capacity, optional)  
logical log volume (transactions, redo-info)  
physical log volume (physical data information)

## 8.4 versant manager

manipulates, caches, provides, marshals objects  
transaction management  
distributes requests for queries, updates to server

## 8.5 object descriptor table (ODT)

logical object identifier (LOIC) → memory, db location  
used on property access, defines if retrieved from memory or db

## 8.6 versant server

updates, caches, retrieves objects  
defines transactions, locks objects  
logging, recovering  
index maintenance

## 8.7 thread architecture

### clients

multiple session with own cache  
have multiple assigned client threads

### servers

multiple server threads process client requests  
access page cache, respect its lock table  
log buffer thread with async IO of uncommitted writes  
background page flusher which writes modified pages

## 8.8 java versant interface (JVI)

store java objects (fine with GC, multithreading)  
client-server architecture (local cache, server queries)  
sits below java VM

### 8.8.1 fundamental layer

database-centric, handlers manipulate objects  
class & attributes builders define classes  
create handles with LOID, new instances  
use handle to get/put values  
fundamental query/result

### examples

```
//schema definition
AttrString name = session.newAttrString("name");
AttrBuilder nameBuilder = session.newAttrBuilder(name);
var attrBuilders = session.withAttrBuilders({ nameBuilder })
ClassHandle person = attrBuilders.defineClass("Person");
person.createIndex(\name", Constants.UNIQUE_BTREE_INDEX);
```

```
//data manipulation
person = session.locateClass("Person");
Handle florian = person.makeObject();
florian.put(name, "Florian");
String val = florian.get(name);
String loid = florian.asString();
florian = session.newHandle(loid);
```

```
//querying
FundQuery query = new FundQuery(session, "select name from Person");
FundQueryResult result = query.execute();
Handle resultSetMember = result.next();
```

### 8.8.2 transparent layer

language-centric, maps classes & attributes to fundamentals  
persistent java object caching & retrieval

### first class object (FCO)

have LOIC, save, query, retrieve individually  
changes saved automatically, applied on commit  
references to other FCO always valid

transient fields not in db  
deleteObject() for db, collected by GC later (then finalize() called)

### second class object (SCO)

saved as part of FCO, can't be queried  
java byte stream if no versant type  
transient fields not in db  
FCO references serialized separately  
changes applied only if owner marked as dirty on commit  
if in two FCO, will be two different instances after fetch  
delete implicitly by reference removal

### persistence categories

for FCO (p, c) possible, marked dirty on modification  
parent class must be same category  
(p) always, new instances directly persistent  
(c) capable, makeRoot(), makePersistent() or reachable persistence  
for SCO (d, a, n) possible  
(d) transparent dirty owner, sets owner as dirty  
(a) persistence aware, can modify FCO, must call dirtyObject() explicitly  
(n) not persistent, can't access fields of persistent object

### persistence model

persistence by reachability  
can elect named roots of graphs for retrieval  
navigate starting at identity, root, class, query  
versant transparently locks & retrieves

### example

```
TransSession session = new TransSession("myDB");
Set<P> ps = session.findRoot("rootName");
Person florian = new Author("Florian");
ps.get(0).addAuthor(florian);
session.commit(); session.endSession();
```

### 8.8.3 ODMG

language-centric, transaction-model, collections  
are FCO, follow ODMG standard

### queries

additional functionality, iff persistent collection, (p) objects  
existsElement, query, select, selectElement  
only elements of collection queried

### VQL

complex expressions, server-side sorting, indexing  
query string which is compiled, optimized, executed on server  
parametrization with \$sign, late binding, can rebind

### example

```
Publication pub = new Publication("Web 2.0 Survey");
String q = "select name from Author where name = $name";
Query query = new Query(session, q);
query.bind("name", "Stefania Leone");
QueryResult result = query.execute();
```

```
q = "select name from Author where Author::Books subset_of $books";
query = new Query(session, q);
query.bind("books", new [] {florian});
```

## 8.9 application development

persistence aware java classes  
deal with sessions, transactions & concurrency  
specify persistence category for classes  
enhancer performs byte-code changes  
create db & run application

### byte code enhancement

create object in db on first instance construct  
read/write objects, attributes to/from db  
FCO pointers evaluated using ODT

## 8.10 sessions

all actions must be performed in sessions  
access to db, methods, data types, persistent objects  
multiple sessions possible  
must close sessions explicitly  
client session has object cache, ODT  
server session has page caches in shared memory

### 8.11 transactions

always in transaction, commit/rollback starts new one  
endSession commits last one  
atomic, consistent, independent, durable, coordinated (with locking), 2PC  
commit() flushes cache, releases locks  
checkpointCommit() retains caches, locks  
commitAndRetain() retains caches, releases locks

### 8.12 object lifecycle

creation of persistent objects (memory, versant cache)  
commit (data written to db, hollow proxy remains in memory)  
rollback (new objects will be dropped)  
query (evaluated on server, proxy created for each result)  
access (fetch or deserialize object)

### 8.13 JVI client cache loader

client-side object cache, server tracks state of clients  
contains query results, navigation results  
dereference consists of RPC, object lookup, IO

#### improve efficiency

vendor specific batch loading

#### configurable strategies

breadth (other trees, same level), depth (deeper level), path loading

### 8.14 collections

standard collections supported (list, array, hashtable)  
FCO collections (VVector, VHashtable)  
SCO collections (DVector)  
FCO large collections (LargeVector, fine-grained locking)  
ODMG collections

### 8.15 event notification

from db to registered clients  
class events (CRUD of any instance)  
object events (CRUD of certain element(s))  
transaction demarcation (begin/end transaction)  
user-defined events

#### event channels

register listeners to channels  
global namespaces of channels over applications  
class, object, query based  
EventClient, ChannelBuilder

#### persistent object hooks

at any sort of state changes  
transient attributes, caches, housekeeping of integrity  
activate, deactivate, pre/post read/write, delete  
can change other objects in hooks

### 8.16 schema evolution

add/rename/remove leaf classes  
change class methods, attributes  
does lazy updates of instances

### 8.17 polymorphic indexes

enhance performance of retrieving object with its subclasses

## 9 social network analysis

### 9.1 social networks

connected people (calls, chats) represented as graphs  
calculates degree, closeness, betweenness centrality

#### find out

key persons of a group of people  
message traversal through group  
communication patterns

#### nodes

key/value pair of people

#### edges

key/value with properties & assigned weight  
uni/bidirectional, explicit/implicit, short/long, single/multiple  
traffic with particular keyword  
work/behaviour patterns

activity such as transfers, payments

### 9.2 graph

#### 9.2.1 path lengths

150 max social relationships (dunbars number)  
4.56 avrg distance of publications connected to erdos (erdos number)  
2.946 avrg distance of movies made with kevin bacon (bacon's number)  
6 avrg distance to know everyone on the world

#### 9.2.2 node properties

##### degree centrality

number of direct connections of a node

##### betweenness centrality

between two important nodes  
high influence over what flows through the network

##### closeness centrality

node with shortest paths to all others  
can monitor information flow best

#### 9.2.3 network structure

##### network centralization

centralized if one or few central nodes  
removing these nodes leads to fragmented network

##### density / cohesion

#direct\_ties / #total.possible

##### distance

minimum number of nodes to connect two specific nodes

##### clustering coefficient

likelihood two associates of certain node are associates too  
high means high clustering

### 9.3 general model

combines different sources, different formats to single model  
uniform analysis, uniform result presentation  
description in triplets (subject,attribute,object) like Linked Data

## 10 graph databases

### 10.1 general

#### meta model

graph containing vertices, edges  
edges, vertices may have key-value properties

#### API

supports CRUD of metamodel  
maybe support traversal of graph  
maybe has graph algorithms implemented

#### characteristics

ACID, scalable for graphs & big data, REST api

#### examples

Objectivity InfiniteGraph, Neo4j, OrientDB

### 10.2 infinite graph

on top of Objectivity, has graph types & algorithms  
distributed graph database

#### usage

extend BaseVertex, BaseEdge to use  
markModified() after modify, fetch() before read  
graphDb.addVertex(v), grapgDB.addEdge(v1, v2,  
EdgeKind.BIDIRECTIONAL)  
graphDb.getNameVertex("name")  
trans = graphDb.beginTransaction(AccessMode.READ.WRITE)

#### navigator engine

result quantifier (append path to results)  
result handler (what happens with path in results)  
path qualifier (continue path or not)  
path guide (which way to continue, DFS or BFS)  
implement Qualifier.qualify(Path p)

### 10.3 Neo4j

#### usage

```

new GraphDatabaseFactory().newEmbeddedDatabase("name")
graphDb.beginTx() tx.success(), tx.failure(), tx.finish()
n = createNode(), n.setProperty("prop", "val")
rela = n.createRelationshipTo(n2, R.MY_TYPE),
rela.setProperty("name", 30)

```

### integrate with java

write wrapper with Node n as private final

### graph traversal

```

td = graphDb().traversalDescription().breadthFirst()
td.relationships(R.MY_TYPE, Direction.outgoing), td.evaluator(eval)
tra = description.traverse(n), for Path p: tra
implement evaluate(Path p) returning exclude/include, continue/prune

```

### algorithms

shortest paths, given length paths, all paths between n1, n2

### rest API

CRUD, traversals, algorithms

### cypher

query language based on pattern matching (START, MATCH, RETURN)  
 START n=node(12) MATCH n-[:author]→b RETURN b.email  
 relationship patterns include (A) → (B), A-[:coauthor]→B

## 11 objectivity DB

### 11.1 general

container-based architecture w/ physical identity

#### meta

OODBMS since 1995, v10, C++  
 java, c#, python, smalltalk, ... frontends  
 data replication, fault tolerance options  
 all platforms like windows, solaris, linux, mac  
 cloud computing in AWS  
 customers from all branches

#### ideal applications

store, process complex structures (trees, collections, graphs)  
 relationship hunting, protein structure, correlation analysis

#### client architecture

languages interfaces provide access to objects & schema  
 local storage / transaction cache  
 client objectivity server (data from local storage, remote processing)

#### server architecture

lock server which grants permissions  
 query server to run queries  
 data server which handles the memory

#### performance

clustering & multi-dimensional indexing  
 client side, cross-transactional caching

#### parallel query engine (PQE)

client side task splitter which can aim queries at specific dbs, containers  
 can split query to multiple agents for parallel processing

#### storage scopes

federation (schema, database catalog) as a file (world.fbd)  
 databases (container dialog) as a distributable file (person.world.DB)  
 container (page map, for logical partitions) consisting of pages

#### storage

each scope/hierarchy can contain up to 2<sup>16</sup> of lower level  
 pages exists as logical, physical, transferred & locked as an unit  
 objects (consisting of slots) stored in container, addressed with page id

#### page map

maps physical to logical pages, journal file saves mapping  
 on transaction, changes persisted to new defragmented page  
 on transaction commit, page map is updated, lock released

### 11.2 c# persistence designer

used to update schema  
 generates c# objects, can generate federation files

### 11.3 persistent object model

basis types ("primitives")  
 complex types, embedded in parent or referenced (OID stored)  
 enumerations, collections, relationships

### 11.4 relationships

unary, binary, to-one, to-many  
 referential integrity maintained by system (incl. inversion)

#### storage

default, non-inline (array stores (identifier, OID))  
 inline (stored as fields on object, to-many stored in array)  
 binary associations as complete separate construct

#### propagation

deletion, locking can be propagated over relations  
 developer specifies how

#### versioning

when object is copied, specify what happens to relations  
 copy (new, old associated with same objects)  
 drop (copy does not have references set)  
 move (copy has the references set, original does not)

### 11.5 domain classes

partial classes from .NET to separate application/persistence code  
 persistence implemented in base class  
 author.cs contains public get;set to private props  
 authorpd generated with private get;set; which do persistence stuff  
 support class w/ schema class, attributes & its properties, proxy cache

### 11.6 connection architecture

static functions to startup(), open connection, shutdown()  
 one connection to federation per application  
 n sessions w/ cache, transaction state, has one, many threads  
 cache kept after commit, flushed on abort  
 if update too big, overflow pages prewritten to disk

#### interaction

connection - > create session → begin transaction, get federation  
 federation will lookup database, create new database

#### persist objects

give reference of db, container, other entity in constructor  
 they are connected, persisted automatically

#### persistent collections

sets, lists, maps as ordered, unordered and scalable, non-scalable variants

#### iterator

provides access to objects meeting certain criteria  
 scope is collection, container, db, federation  
 criteria is PQL predicate as a string  
 not efficient unless indexes preconstructed  
 lazy filtering, therefore no sorting

#### scope name

can name object, collection at each storage hierarchy (like roots)

#### retrieve objects

by scope name, link following, lookup with keys & iterators  
 parallel query with PQE  
 content-based filtering (supporting primitive types, group lookups)

#### retrieve objects by group

use object iterator for storage hierarchy, name maps, root names, name  
 scope  
 use collection iterator for lists, sets, object maps

#### LINQ

language integrated queries, transforms query in method calls  
 objects (in memory), SQL (MSSQL), XML, DataSet (ADO.NET)  
 other providers possible like to db4o

## 12 Object Store

page-based architecture w/ physical identity  
 query may be executed on server

### 12.1 meta

#### personal edition

lightweight object database  
 large, single user database  
 multithreading, small memory footprint  
 for embedded systems, mobile computing, desktop applications

#### enterprise

distributed multi-user

object caching  
for clustering, online backup, replication, high availability

## 12.2 virtual memory mapping architecture VMA

extends OS memory management for persistence  
data is referenced with (database, segment, cluster, offset)  
translated to place in virtual memory (the PSR, see below)

### persistent storage region (PSR)

PSR sits between heap & stack  
acts as new layer of cache, data of client inside this area  
serves as secondary storage, persistent over transactions  
implemented with indirect pointer

### page fault

if data accessed in PSR which is not there yet  
ObjectStore maps from memory or server to location

### address translations

done when data is fetched into cache  
retranslation may occurs when PSR full

## 12.3 server side architecture

### server

manages databases & transaction logs  
serves content in pages  
enforces ACID with permits  
2PC with other servers  
recovery mechanisms

### database

binary files storing pages of c++ memory

### transaction log

pending changes of transactions, on commit executed on db  
for recovery, faster commits, multi-version concurrency control MVCC

## 12.4 client side architecture

### client

C++ using ObjectStore library  
pages automatically fetched from db according to demand

### cache

one cache memory file per process with fixed size  
contains pages fetched from the server (over transactions)

### commseg

one commseg memory file per process describing the cache  
stores permit, lock, meta data for all pages (kept between transactions)

### cache manager

one process by client, shared by all clients on machine  
handles permit revokes, read/write to cache & commseg

### persistent storage region PSR

reserved area between heap/stack  
logical pointer addresses of client map into this area  
ranges with fetched pages, ranges with to-be-fetches pages  
if (PSR full) then updated pages back to DB, read-only dropped  
if (end of transaction) then PSR cleared to be reused

## 12.5 fetching/mapping pages

client fetches automatically, lazily  
server permits & client locks acquired automatically

### steps

ObjectStore override segment violation SIGSEGV handler  
program obtains pointer p to page x  
access to x causes exception, caught by SIGSEGV handler  
page fetched from server, put into cache  
execution continues

## 12.6 cache forward architecture (CFA)

permits are tracked by the server (represent ownership)  
locks taken by client according to permit (no-lock, read, write)  
on page request, client checks for permit states & resolves conflicts

### locks

read permits to take readlock (multiple clients simultaneously)  
write permits to take writelock (only single client)  
else client must ask server to escalate permissions

### server call-backs

to revoke read permit (if server needs write)  
to revoke write permit to read  
ACK response if client has no lock taken  
NACK response if client has lock, but needs to give up after transaction

### allows high performance

data cached across transactions  
fewer locks need to acquired  
cached data in globally consistent state

## 12.7 distribution & heterogeneity

transactions can include objects from different db  
different client platforms as runtime marshals

## 12.8 persistence

persistence by instantiation C++  
orthogonal to types  
single db can be used for both transient, persistent objects

### allocation

overloaded new(allocation, type specification, count = 1)  
locate transiently on heap  
locate on database, segment, cluster, next to other object

## 12.9 transactions (ACID)

atomicity (commit saves changes, abort removes changes)  
consistency (impossible to apply, lose data on write)  
isolation (2PL serialisability, MVCC snapshots for read-only transactions)  
durability (changes to transaction log, background processes write to db)

### types

read (throws exception if write requested)  
write  
local (initiating thread can execute)  
global (all threads in session can execute)  
lexical (thread-local, start-end in same code block, automatic retry)  
dynamic (lower level, for multi-threaded application)

## 12.10 database layout

clusters hold segments, saved as pages

### segments

logical partitioning of object  
0 for schema, db roots  
2 default  
4+ user created till  $2^{32}$

### clusters

group closely related objects  
0 default cluster  
1+ user created

## 12.11 developing applications

### 12.11.1 basic constructs

#### objectstore

runtime of ObjectStore  
static initialize(), static shutdown()

#### os\_database

database functionality, applies updates automatically  
static create(), save(), close() (does not save state), destroy() (deletes)

#### os\_transactions

transaction handles & functions, can be nested arbitrarily  
static initialize(), static get\_current(), use mixins block pre/postfix  
abort(), commit()

#### os\_typespec, os\_ts

determine type specification  
Author\* = new(db, os\_ts<Author>::get(), 1) Author("Stein")

#### os\_database\_roots

creation, retrieval, removal of roots

#### os\_segment

segment management

#### os\_cluster

cluster management  
static of(object\_reference) returns cluster of that objects



## os\_database\_root

labelled persistent objects

create\_root(), find\_root(), get\_value()

### 12.11.2 development process

writing of persistent classes, schema file, application logic

schema file compilations with pssg compiler

C++ compilation, then link with pssg binaries

### 12.12 relationships

reference in both objects using os\_collections

### 12.13 os\_collections

fit for traversal, manipulation, retrieval

set, bag, list, array, dictionary, also available as generics (templates)

new(db, os\_Set<T>::get\_os\_typespec()), insert(), delete keyword

reusable cursor with T first(), T next(), boolean more()

#### queries

specify element type, query string, database

query can be c++ condition as string or regex

nested queries & basic function calls allowed

## 13 the OM data model

### 13.1 elements

multiple inheritance, instantiation, classification

collections, associations

cardinality, classification, evolution constraints

### 13.2 about

extended entity-relationship model for OO data management

distinguishes typing, classification (representation vs roles)

data represented as objects (attributes, methods)

multiple inheritance, instantiation, classification

collections (set of persons), binary collections (person ↔ location)

integrity, classification, evolution constraints

data definition, manipulation & query language OML

### 13.3 typing & classification

clears up issues by separating the two concepts

less complex type graphs but still rich classification possible

#### typing

employee extends person (naive) vs collection concept (semantic grouping)

for representation, data format, operations, inheritance

#### classification

person, employee → semantic grouping → contact, address, name, jobinfos

for roles, semantic grouping & constrains (member types) in collections

to define relationships in collections (no embedding in objects)

to add/remove roles dynamically

### 13.4 vs state machines

view collection membership as state

assign methods when collection membership changes

react on events

### 13.5 OM data model layers

#### 13.5.1 type layer

objects represented by multiple object types

object types define type units

dress operation to add type, strip to remove (dynamically)

multiple inheritance (multiple supertypes, developer handles conflicts)

multiple instantiation (can have unrelated types from different hierarchies)

#### supported values

base types without identity (string, int, uri)

object types (objects with identity)

structured types (structures without identity)

bulk types (collection of same member, no identity) (arrays)

#### type units

corresponds to type, no inherited fields, methods (traits)

has attributes with name, type, bulk (name, string, uni)

#### information unit

instances of type unit has corresponding information unit

contains the data of the type unit

browse information unit by casting to its type given object identifier

person(o01) = ("max"), private(o01) = (1234, Bern)

#### bulk types

uni (single), set, bag, sequence, ranking

#### dress, strip

can dynamically add/remove type instances

dress create information unit with default units and attaches

strip removes information unit and discards values

### 13.5.2 classification layer

defined based on types from type layer

defines semantic grouping, multiple classification, collections, associations

#### collections

concept for semantic grouping, associations to link objects together

membership constrained by type (multiple collections can have same

constraints)

kinds & roles

ext(Persons) = { person(o1), person(o2) }

#### collection behaviours

set Person (no duplicates, no order)

bag <Person> (duplicates, no order)

sequence [Person] (duplicates, order)

ranking [Person] (no duplicates, order)

#### subcollection behaviours (denoted with arrows)

arrow from collection to other collection

equal (same elements, maybe different behaviour / types)

strict (subsequence, subranking)

total (all instances of type from larger collection)

#### structure

arrows from multiple source collections to single target collection

disjoint (at most in one source collection)

cover (in at least one source collection)

partition (in single source collection)

intersection (in all source ⇒ in target, arrows other way)

#### associations

source, domain collection ⇒ relation collection ⇒ target, range collection

relation collection is of the form (source(o1), target(o2))

cardinality constraints (0:\* means ≥0, 1:\* means >0)

behaviour represented same as normal collections (set, bag, etc)

ternary, attributed relations not supported

#### nested associations

domain, range collection can be relations too (for n-ary relationships)

ternary relationships decomposed into primary, secondary

allows uniform query constructs, clearer semantics of relations

#### kinds, roles

kind (fixed not-changable classification, person; professor; postgrad)

roles (change during entity lifecycle, baby → adult → senoir; student)

#### constraints, structures

applicable to subcollection, superconnections

for cardinality (describe associations)

for evolution (govern object lifecycle)

#### evolution

assume each collection has single root (is a kind), single maximal collection

kind is fixed classification in context of parent (senior of kind human)

therefore kind can only be striped if parent can be striped too

#### determine valid evolution

migrate element x from C\_1 to C\_2, denoted as x :: Postgrad → Lecturer

(1) x does not belong to subcollection of C\_1 (trivial if C\_1 is leaf node)

(2) let K be kinds(C\_1) - kinds(C\_2) then all roles(K) can't be in C\_2

#### example valid evolution

x :: Postgrad → Lecturer

(1) trivial because Postgrad is leaf node

(2) K = {Postgrad, Person} - {Person}, no r ∈ roles(Postgrad) in

roles(Lecturer)

### 13.6 object model language (OML)

declarative, object-oriented, for object data model

#### OML data definition language

object, structured type definition

method definition, implementation

collection, association, constraint definition

## OML data manipulation language

CRUD, dress & strip operations

## OML query language

expressions, functions for base type values

operations to access properties, execute methods for objects

operations on collections based on collectional algebra

### 13.7 schema definition

```
create type contact { name : string, webpage : uri }
create type person subtype of contact { age : int }
method getWork() returns (location: set of location) ( )
```

### 13.8 classification definition

```
create collection Contacts as set of contact
create collection WorksFor as set of (person, organisation)
create constraint con1 association on WorksFor from Person(0,*) to
Organisation(0,*)
create constraint con2 subcollection Person restricts Contacts
create constraint con3 classification (Person, Organisation) partition
Contacts
create constraint con4 classification Person is kind
```

### 13.9 queries

applies to collections, binary collections

#### selection

retrieve person object matching predicate  
all \$p in persons having (\$p.title = "prof")

#### map

retrieve table with attributes as columns  
map \$p in persons by (\$p.title x \$p.name)

#### extraction

specify first/last, top nth elements, max/min  
the 3 in persons; first persons, max persons.getAge(),years

#### reduce

calculate aggregated value  
reduce \$p in persons aggregate \$a by (\$p.age + \$a) default 0

#### domain, range (binary)

get left (domain), right (range) part of binary collection  
domain locatedAt, range locatedAt

#### domain/range restriction (binary)

retrieve associations which match predicate  
locatedAt dr (all \$p in persons having (p\$.title = "prof"))

#### domain/range subtraction (binary)

contrary of restriction (takes other results)  
locatedAt rr (all \$p in persons having (p\$.title = "prof"))

#### inverse (binary)

swap domain/range of collection  
inverse locatedAt

#### nest (binary)

groups range of each domain in set  
nest locatedAt

#### compose (multiple binary)

join range of first with domain of second  
situatedAt compose inCountry

#### closure (binary)

compose binary with itself (for hierarchic relations)  
closure partOf

#### division (binary)

divide binary by collection, keep domain/range in all results  
situatedAt div cities

### 13.10 collection algebra

union U, intersection N, difference -, selection %, map alpha, reduce o+,  
flatten +-  
extend standard collection algebra for bags, sequences, rankings

#### bags

as tuples,  $\langle a, a, b \rangle \rightarrow \{(a, 2), (b, 2)\}$   
union by taking  $\max(n_1, n_2)$  of both sets  $\langle a, a \rangle \cup \langle a, b \rangle = \{(a, 2), (b, 1)\}$   
addition by taking  $\text{sum}(n_1, n_2)$   
intersection by taking  $\min(n_1, n_2)$   
 $B_1 - B_2$  by taking  $n_1$  if not in  $B_2$ , or  $n_1 - n_2$  if in both

$B \% P = \{(x, n) \mid (x, n) \in B \wedge P(x) = \text{true}\}$  for  $P(x) = x \rightarrow \text{bool}$   
 $B \text{ alpha } f = o+ B((x, n), B_1) \rightarrow (\{(f(x), n) \mid (x, n) \in B\})$  for  $f(x) = x \rightarrow y$   
 $o+ B \text{ f } v = f(x, o+ B' \text{ f } v)$  where  $B = B' \text{ union } \{(x, 1)\}$  for  $f(x) = x_1, x_2 \rightarrow y$   
 $+ - B = o+ B(x, B_1) \rightarrow (\text{ext}(x) \cup B_1)$  (flattens 3D to 2D)

## 14 storage and indexing

type hierarchy indexing

aggregation path indexing

collection operations

### 14.1 motivation

manage large data set on persistent storage

#### more requirements than relational

structuring, clustering, management of complex objects

access through references, query predicates

type inheritance hierarchies

relationships

multi-values properties & collections

### 14.2 object oriented storage

layouts not very different from relational systems

#### new algorithms

data structures for complex objects

grouping/clustering of complex objects

grouping/clustering of references

management of free space / buffer

### 14.3 storage model

#### value structures

consisting of data sets (records, attributes, domains)

divided in pages (disk block, sequence of disk blocks)

functions to map records to pages

#### access structures

query/update algorithms

search data structures

### 14.4 terminology

#### query

point, range query (exact value, interval matching)

single class / hierarchy (instances with / without hierarchy)

#### index

unique, non-unique (key, non-key fields)

sequential, non-sequential keys (ordered, unordered)

one / multi-dimensional (index over single, multiple fields)

compound (one-dimensional index over multiple concatenated values)

placing/clustering (search physical structure)

#### sequential organisation

pages organised by data set, new entries are placed into newest page

queries have to traverse entire data set

data index maps ids to physical locations

#### subspace mapping

data set decomposed into subspaces (can overlap)

queries traverses B-trees, K-trees, grid files (x,y boundaries)

leaves contain physical location of entries

#### point mapping

data set directly mapped to specific place in memory

query by using hash functions

directory with all points references linked pages

### 14.5 data structures

#### B-tree

each block contains n indexes, n+1 references

balanced,  $O(\log n)$

#### B+-tree

like B-tree, but leaves contains direct value

redundant indexes pulled down for references to the right

easier for aggregated searches, breadth-first search

#### linear hashing

bucket for each hash result

### **extensible hashing**

start with smallest index as possible (e.g. 2bit keys)  
extend if the need is up for it

### **bloom filter**

m sized array with boolean entries, n hash functions  
entry n times hashed, activates target boolean entries  
for fast existence test

### **inverted files**

search structure called vocabulary (indexes distinctive keywords)  
inverted list for each keyword, storing id of records with that keyword  
sort inverted list to apply compression

### **signature files**

each record has a fixed width index information  
keywords are hashed, and looked up if index information indicates  
use bloom filter

### **k-d tree**

tree with k dimensions indexed, splits region into subregions  
search for key with multiple dimensions closest match in tree

### **r-tree**

similar purpose than k-d tree  
differences include balanced, disk-oriented, rectangle partitioning

### **H-tree**

fractal built up tree

### **others**

hB-tree, Quadtree, TV-tree, cell tree

## **14.6 type hierarchy indexing**

when using object in query implicitly using type hierarchy  
build index with type or key as top-level element

### **single class index (SC-index)**

construct search structure for each type containing its subtypes  
query evaluator needs to traverse for all used components

### **class hierarchy index (CH-index)**

one search structure for all indexes types  
query evaluator scan through once, collects all types

### **h-tree**

skipped, explanation not clear at all

### **class division index (CD-index)**

compromise between query and storing all in one node  
q (# of search structures for any type), r (# of replication of types)  
types are stored aggregated in a node, or combine freely at runtime

### **multi-key type index (MT-index)**

type membership is just another attribute  
on evaluation, collect disk addresses and disregard if not qualifying

## **14.7 aggregation path indexing**

avoid full traversal of paths & intermediate object loading  
nested index provides direct access between start/end objects of relation  
path index stores all paths to ending objects (predicates possible)

### **multi-index (MX)**

divide paths into subpaths of length one  
for each relation, stores source/target relation  
for backwards traversal (start at attribute, then get objects)

### **access support relations (ASR)**

can-extension aggregates paths from 0 to n  
left-complete extensions with all paths from 0, possibly till n  
right-complete extension with all paths to n, possibly starting at 0  
full extensions with all paths, possibly starting/ending at 0/n  
answer queries starting at end points efficiently

### **nested index (NX)**

backwards traversal of the full path  
equivalent to backwards traversal of ASR can-extension

### **path index (PX)**

backwards traversal of right-complete ASR

### **join index (JX)**

keep binary join indexes at both sides of the relation

## **14.8 collection operations**

multi-valued attributed evaluation

### **signature files (like bloom filter)**

element signature summed up equal signature for object  
create signature for query too, then compare with object signatures

## **15 version models**

### **15.1 domains**

temporal databases  
computer aided design & manufacturing  
software configuration / engineering

### **15.2 basic aspects**

#### **granularity**

files, tuples of relation, attributes of class, entire object

#### **organisations**

set, list, tree, (directed acyclic graph) DAG

#### **reference types**

specific (reference single version of object)  
generic (references object, upon usage dereference specific version)

#### **storage**

complete version of objects  
forward/backward state/operation-based delta (changes) between versions  
space vs performance tradeoff

#### **operations**

create / branch / merge / delete  
similar to transactions (long running, nested)

#### **interaction models**

automatic versioning  
explicit high level user operations (check-out, commit, update)  
implicit query expansions

#### **queries**

active versions guide to query sequential versions  
main derivation guides to query parallel versions  
may combine for generic reference

### **15.3 temporal databases**

first application of version models

#### **notion of time**

AS-OF operation (time of transaction, physical time)  
WHEN operation (real world occurrence)  
user defined time

#### **classifications**

static/snapshot (single version managed)  
static roll-back (AS-OF, space overhead)  
historical (WHEN, can change historic data)  
temporal (both WHEN / AS-OF)

#### **spatio-temporal data**

moving objects db, position uncertainty  
(sometime/always,possibly/definitely) inside (trajectory, range, t.1, t.2)  
therefore 8 possible operators, some semantically equivalent

#### **representations**

tuple versioning (keep old rows, new column defines newest version)  
attribute versioning (each attribute has temporal info)

#### **bitemporal conceptual data model (BCDM)**

adds columns transaction time, valid time, until changed, now (boolean)  
TSQGL 2 with VALIDTIME, WHEN clauses

#### **models**

homogeneous if all attributes changed at the same time  
heterogeneous if attributed changed at different times

#### **anomalies**

vertical if multiple tuples for single entity (tuple versioning)  
horizontal if entity spread over multiple relations (attribute times vary)

#### **storage models**

primary store for current versions, for non-temporal queries  
history store for other queries  
reverse chaining (references to next older version)  
accession lists (reference to version table, then version access)  
stacked versions (combine accession lists & reverse chaining)

## 15.4 engineering databases

for computer aided design/manufacturing CAD/CAM  
for dev/maintenance of objects

### requirements

handle complex, hierarchical object structures  
support versioning through incremental development / trial and error

### dimensions

linear revisions  
(non-sequential) variations

### modelling primitives

component hierarchies (is-part-of)  
version histories (is-derived-from)

### design management

identify current version, describe dynamic configuration  
change & constraint propagation (inherit attributes from related)

## 15.5 software configuration systems

goal to fully automate building final product  
build around design objects (source code, modules)  
manages dependencies & references

### product space

describes product organisation  
several design choices possible, e.g. software with different requirements  
objects can be specified at different granularities, representations  
relationships (root is product, visualizes dependencies)  
logical structure (modules import others explicitly)  
file system (module in files, build file defines relationships)  
data model (tree with dependencies as leaves, generate build info)

### version space

defines how objects are versioned  
keeps invariants and deltas  
revisions keep track of history, variants capture alternatives  
one-level representation (sequence, tree, DAG)  
two-level representation (revisions, variants)