

# Advanced Systems Lab - FFT

8337 characters in 1374 words on 250 lines

Florian Moser

July 8, 2020

## 1 fast fourier transform (FFT)

### 1.1 introduction

#### history

1805 Gauss discovered FFT for personal paper-interpolation  
signal processing done analog (with currents)  
1965 FFT cooley tukey rediscovers; enables digital processing

#### purpose

change of basis by multiplying with fixed matrix  
is a linear transform (n-vector input/output)  
 $y = Tx$  ( $y_k = \sum_{l=1}^n t_{kl} x_l$ )

#### optimization potential

up to 35x faster  
5x for locality, 3x for vectorization, 3x for threading

### 1.2 discrete fourier transform (DFT)

$y_k = \sum_{l=1}^n (e^{-2\pi i k l / n}) x_l$   
all n roots of unity (of one)  
 $w_n$  called primitive nth root of 1

#### notation

$y = DFT_n x$ ,  $2n^2$  operations

#### evaluate

get multiplier using n (size)  $m = kl \cdot 2\pi / n$   
insert  $k/l$  (row/column; zero-based)  
then evaluate with  $\cos(a) + i \sin(m)$   
remember that  $\sin(0) = 0$ ,  $\cos(0) = 1$

#### $DFT_2$

$\begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$   
no mults (bc all factors 1), 1 adds/sub  
 $y_1 = x_1 + x_2$ ;  $y_2 = x_1 - x_2$   
called butterfly

#### $DFT_4$

$\begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & i & -1 & -i \\ 1 & -1 & 1 & -1 \\ 1 & -i & -1 & i \end{bmatrix}$   
12 adds, 4 mults (if not reduced)  
in complex arithmetic, needs to be mapped to machine

#### further examples

DFT & RDFT universally used  
MPEG & JPEG uses own kind of DFT

### 1.3 cooley-tukey FFT

FFT transformation algorithm  
calculate T like  $T = T_1 * T_2 * \dots * T_m$   
only useful if  $T_i$  are sparse, and m low

#### 1.3.1 FFT=4

essentially 4 DFT of size 2

#### matrix

for  $n = 4$  (better readability)  
 $T_1 = [1, \dots, 1, \dots, 1]$  (shift; 0 ops)  
 $T_2 = [1, 1, \dots, 1, \dots, 1, \dots, 1, \dots, 1]$  (four adds)  
 $T_3 = [1, \dots, 1, \dots, 1, \dots, 1, \dots, 1]$  (one mul by i)  
 $T_4 = [1, 1, \dots, 1, 1, \dots, 1, \dots, 1, \dots, 1]$  (four adds)  
 $y = T_4 * T_3 * T_2 * T_1 * x$   
 $\Rightarrow$  reduced FFT opcount to 8 adds, 1 mul

#### algebra

$(DFT_2(x) L_2) \text{diag}(1, 1, i, i) (L_2(x) DFT_2) L_2^{-1}$   
 $I_2 = \text{diag}(1, 1)$  is identity matrix of size 2  
 $L_2^{-1}$  = permutation with stride 2  
for (x) kroenecker product  
multiply entry on left with whole matrix right  
produces for  $2 \times 2$  (x)  $n \times n = 2n \times 2n$  result

#### dataflow representation

right  $x_1, x_2, x_3, x_4$  flows to left  $y_1 y_2 y_3 y_4$   
permutation of  $x_2$  and  $x_3$  ( $L_2^{-1}$ )  
 $DFT_2$  applied to  $x_1, x_2$  and  $x_3, x_4$  ( $L_2(x) DFT_2$ )  
scaling represented as black dots ( $\text{diag}(1, 1, i, i)$ )  
 $DFT_2$  applied to  $x_1, x_3$  and  $x_2, x_4$  ( $DFT_2(x) L_2$ )

#### 1.3.2 generalization

$(DFT_k(x) L_m) T_m^{-1} (L_k(x) DFT_m) L_k^{-1}$   
choose radix k, select m to fit input  
 $DFT_k(x)$  results with  $DFT_k$  in the diagonal  
 $T_m^{-1}$  as diagonal with m-1 1s, then an i, then repeat  
 $L_k(x) DFT_m$  results in m A's at stride m  
 $L_k^{-1}$  permutation; read at stride k, write at stride 1  
can reformulate to read at stride 1, write at stride m  
get n logn algorithm

#### variants

decimation-in-time  
decimation-in-frequency (transposed cooley-turkey)

#### factors

can choose radix k, but in the end  $k \cdot m = n$   
for factors of 2 trivial; for primes need different algorithm

#### cost

$n \log(n)$  complex adds; each needs 2 real adds  
 $n \log(n)/2$  complex mults; each needs 2 adds, 4 mults  
hence  $3n \log(n)$  adds,  $2n \log(n)$  mults =  $5n \log(n)$   
reduce op count depending on recursion  
because can simplify  $*i$  or  $*1$

#### recursive vs iterative

calculation order different (DAG equivalent)  
iterative in stages (compute all of same kind of butterfly)  
recursive in recursion (compute butterfly as soon as dependency resolved)  
only order of operation changes, not count  
recursive better on caches bc temporal locality

### 1.4 FFT variants

#### iterative

initial permutation step  
then bigger and bigger butterflies  
size of butterflies defined by radix

#### pease

permute result of butterflies  
to only operate on 2butterflies; good for CPUs

#### stockham

big butterflies and then permute  
permutations always in pairs of two  
good for 2-element cache lines

#### six-step FFT

optimized for parallelization  
three stages of communication, two parallel stages  
works for all sizes, any computation

#### multi-core FFT

similar to six-steps, but always two elements travel together  
well if block size 2; no false sharing  
hence two processes access elements in same cache line

### 1.5 complexity

use  $L_c$  as measure; defines cost add/mul = 1 for number  $< |c|$   
used  $L_2$  in proofs;  $L_\infty$  would be more real-worldly

#### upper bounds

for  $n = 2^k \Rightarrow 3/2 n \log(n)$   
other  $n \Rightarrow 8 n \log(n)$

## lower bound

$L_c(\text{DFT}) \geq 1/2 n \log_c(n)$

hence as  $n$  goes to infinity then constant  $c$  also goes to infinity

fastest algorithm found is close under  $4n \log(n)$  for  $n=2^k$

fastest algorithm in use needs little more than  $4n \log n$

## 1.6 optimizations

### 1.6.1 choice of algorithm

iterative used to be the best choice

bc easy to implement, fewer instructions

recursive now better

bc caches (temporal locality)

### 1.6.2 locality improvement

trivially need 4 data passes (bc 4 steps)

#### DFTrec

fuse T1,T2 (read stride  $k$ , DFT<sub>m</sub>, write stride 1)

introduce interface DFTrec( $m, x, y, k, j$ )

for  $m$  size,  $x/y$  vectors,  $k$  input stride,  $j$  output stride

out-of-place (reads  $x$ , writes  $y$ )

called recursively (just change stride)

#### DFTscaled

fuse T3,T4 (read stride  $m$ , prescale, DFT<sub>k</sub>, write stride  $m$ )

introduce interface DFTscaled( $k, y, d, m$ )

for  $k$  size,  $y$  input/output vector,  $d$  prescale,  $m$  input/output stride

in-place (reads & writes  $y$ )

base case (hence rewritten for each radix factor)

#### pseudocode

for  $n$  size,  $x/y$  input/output vector,  $t$  prescale lookup

choose  $k$  (depends on  $n$ ; ensure base case for DFTscaled exists)

let  $m = n/k$

for ( $i$  until  $k$ ) DFTrec( $m, x+i, y + m*i, k, 1$ )

for ( $i$  until  $m$ ) DFTscaled( $k, y+i, t[i], m$ )

### 1.6.3 constants

FFT needs multiplications by roots of unity

expensive because needs sin/con to compute

#### precompute

DFT\_init( $n$ ) for  $n$  input size to create lookup table

then reuse result for all same-sized DFT computations

prestore for small input sizes

### 1.6.4 optimized basic blocks

#### unroll recursions

bc hard for compiler to do it automatically

function calls overhead negligible (bc base case reached fast)

#### base cases

to optimize specifically for small inputs

empirically useful for  $n \leq 32$

but then need 31 DFTrec & DFTscaled each

can be generated (as its done by FFTW)

### 1.6.5 adaptivity

adapt recursion strategy to platform

do search within init function (reuse DFT\_init( $n$ ))

use dynamic programming (saves perf. compared to exhaustive)

#### valid recursions

generated droplet for base cases must exist

DFTscaled as a basecase (hence always right-recursive)

## 1.7 FFTW (fastest fourier transform in the west)

generates codelets using SIMD

with input  $n$  generates DFTrec, DFTscaled in three steps

### 1.7.1 DAG generator

create trees based on sum formula

includes cooley-tukey, split-radix, ....

fuse trees into DAG

### 1.7.2 simplifier

simplify DAG

#### algebraic transformations

mults ( $0*x \Rightarrow 0$ , similar cases for  $-1,1$ )

distribution laws ( $kx + ky \Rightarrow k(x+y)$ )

canonicalization ( $x-y, y-x \Rightarrow x-y, -(y-x)$ )

## common subexpression elimination (CSE)

use temporary variables to avoid double computation

## reduce constants (DFT specific)

has many multiplications with many different constants

each constant used positive & negative (bc  $\sin(a) = -\sin(a)$ )

reduce register pressure by storing only positive constant

and use subtraction where  $-\sin(a)$  would be needed

### 1.7.3 scheduler

transform DAG to  $c$  under register spill minimization

want to reach target  $I = O(\log(C))$  for cache size  $C$

uses SSA style, scoping

#### approach

cut DAG in middle

then recurse on connected components (butterflies) to outside

hence use minimal number of registers for "durchstich"

#### cut DAG middle

start from the outsides (left = input, right = output)

color immediate next inner reachable nodes

repeat recursively until nodes touch in the middle

## 1.8 comparison

MMMM (ATLAS)

sparse MVM (Sparsity, Bebop)

DFT (FFTW)

#### cache optimizations

blocking for MMM, SMVM

recursive FFT & fusion of steps for FFTW

#### register optimizations

blocking for MMM, SMVM

schedule small FFT for FFT

#### optimized basic blocks

unrolling

scalar replacement & SSA

scheduling & simplifications for FFT

#### other optimizations

precompute constants for small DFT

for large FFT the algorithm becomes memory bound

hence avoid precomputation to save on storage

#### adaptivity

search blocking parameters for MMM

search register block size for SMVM

search recursion strategy for DFT

## 1.9 spiral

specifically to optimize FFT

for specific input size

applicable for locality, threading, ...

#### algorithm

use signal processing language (SPL)

is mathematical, declarative, point-free (any input size)

divide and conquer applied to transform algorithms

#### generate code

decompose DFT with SPL rules to algorithm

optimize for parallelization/vectorization

transform to structure with sums

optimize for locality

transform to C program

optimize in basic blocks

#### vectorization optimization

some SPL expressions directly relate to SIMD code

like  $DFT_2(x) I_4$  which can be transformed easily

hence rewrite given SPL to vectorizable SPL expressions

#### generate base cases

express intrinsics as matrices

then search for matrixes matching base cases

#### generalize for any input size

need many more codelets

results in huge library size

but faster than other approaches; including FFTW