

# Advanced Systems Lab

30615 characters in 4944 words on 894 lines

Florian Moser

July 7, 2020

## 1 introduction

to improve performance need techniques throughout the stack

### 1.1 motivation

computing needs unlimited performance  
many different algorithms  
only small set of critical components (100s)

#### applications

scientific computing (climate/financial/molecular modeling)  
consumer computing (compression, security)  
embedded computing (sensors, robotics)  
research (biomedics, imaging, computer vision)

#### critical functions

transforms (fourier, etc.)  
matrix multiplication  
filters/correlations/convolutions  
equation solving  
dense/sparse linear algebra  
coders/decoders  
graph algorithms

### 1.2 archive speedups

with same compiler, flags, operation count  
archive 10x - 100x speedup

#### levels

algorithms (lower complexity)  
software (manually optimize low level)  
compiler (automatically improve low level)  
microarchitecture (optimize for specific processor)

#### low level optimizations

adapt to memory layout  
use vector instructions  
use parallelism

#### achievable

20x through memory hierarchy (caches)  
4x through vector instructions  
4x through parallelization

#### optimized code in practice

heavy LoC count, manual process with some generated parts  
often not portable to other microarchitecture  
may violate software engineering practices  
uses vector instructions and parallelization  
reduces cache misses (potentially at the cost of more operations)

### 1.3 naive matrix multiplication (NMM)

for  $n$  times  $n$  matrixes  
 $O(n^3)$  operations  
 $O(n^2)$  space

#### for optimized algorithms

vector instructions reduce number of instructions  
but memory access stays asymptotic same  
because just constants improved

#### for parallelized algorithms

may take  $p$  processors into account  
like  $O(n^3/p)$

#### limitations

asymptotic runtime only describes eventual trend  
constants matter; unclear when "eventual" starts

## 2 basics

### 2.1 o-definitions

$O(f(n))$  is set of functions;  $f \in O(g)$   
 $O(g)$  (only reasonable faster growth)  
 $\theta(g)$  (equally fast growth)  
 $\omega(g)$  (only reasonable slower growth)

#### order

$1 < \log \log(n) < \log(n) < \sqrt{n} < n$   
 $n < n \log(n) < n^{1+e} < n^2 < n^m < 2^n < n!$

#### simplifications

$\log_2(n) \rightarrow \log_3(n)$  (base of logarithm irrelevant)  
 $O(n + \log(n)) \rightarrow O(n)$  (take highest factor)  
 $O(100n) \rightarrow O(n)$  (remove constants)  
 $O(n^{1.1} + n \log(n)) \rightarrow O(n^{1.1})$   
 $2n + O(\log(n))$  (keep, bc. outside bigger)  
 $O(n) + \log(n) \rightarrow O(n)$  (only  $O(n)$  relevant)  
 $O(n^2 + m)$  (keep, bc. different numbers)

### 2.2 cost analysis

count number of relevant operations

#### cost measure $C(n)$

define what cost is composed of  
for example  $C(n) = (\text{adds}(n), \text{mults}(n))$   
only account for mathematical operations  
ignore computer mapping operations  
like array access, index calculations

#### index operations

ignore if simple or dominated by expensive op  
not the bottleneck in good processor (bc very common)  
for example integer & floating point units separate  
 $c[i*m + k] * c[i*n + f]$  (ok; dominated by mul)  
 $c[i + 5k + 4m*m]$  (not ok; complex)  
 $c[i/m \% k] + c[n]$  (not ok; + does not dominate / or %)

### 2.3 performance

equals cost/runtime

#### machine peak performance

max flops possible on that machine  
useful to compare to archived flops  
90% peak percentage possible for matrix multiplication  
2 flops/cycle = 1 fma  
32 Gflops = 4 cores \* 2 AVX \* 1 add \* 3 GHz

#### performance NMM

$C(n) = (\text{fladds}(n), \text{flmults}(n)) = (n^3, n^3)$   
for  $n = 10$ , cost = 2000, runtime = 1000, peak performance = 2  
gives performance = 0.5 hence 25% peak performance

#### use performance counters

count cache misses, floating point operations  
heavily OS/processor dependent  
may not be available / faulty  
like Intel PCM, Intel Vtune, Perfmon

#### instrument code manually

measure time at a specific place in code

### 2.4 operational intensity

$I(n) = W(n) / Q(n)$   
for flops (floating point operations)  $W(n)$   
for bytes transferred (assuming empty cache)  $Q(n)$

#### intensity factors

computer (memory bandwidth, peak performance)

implementation (degree of locality, SIMD)  
measurement (warm cache, cache source, ...)

#### **compute-bound**

if  $I > \text{peak performance} / \text{memory bandwidth}$  ("high", often n factors)  
computation uses same value often; work package in cache  
cache miss time negligible compared to computation  
improve performance by keeping floating point units busy  
use ILP & vectorization  
80-90% of peak performance archivable in practice

#### **memory bound**

if  $I < \text{peak performance} / \text{memory bandwidth}$  ("low", constant)  
computation needs many values; work package too big for cache  
cache miss time bigger factor than computation time  
improve performance by reducing memory traffic  
reduce cache misses & register spills, use compression  
40-50% of peak performance achievable in practice

### **2.5 operational intensity examples**

$y = y + x \quad O(n) / O(n)$   
 $y = Ax \quad O(n^2) / O(n^2)$   
FFT  $O(n \cdot \log n) / O(n)$   
 $C = AB + C \quad O(n^3) / O(n^2)$

#### **fma two array sum**

$W(n) = 2n$  flops needed (add + mul per row)  
 $Q(n) \geq 2n$  (need 2 n-arrays)  
 $I \leq 1/8$ ; hence memory bound  
 $W' = n/2$  without vector ops (use 2 fma ports)  
 $W'' = n/8$  with vector ops (do 4 at same time)  
 $Q' = n/4$  L1 & L2 (do 8 loads at same time)  
 $Q'' = n/2$  L3 (do 4 loads at same time)  
for double, multiply Q by 8 bc double needs 8 bytes

#### **matrix multiplication**

$W(n) = 2(n^3)$  flops needed (add + mul per row/column combination)  
 $Q(n) \geq 3(n^2)$  (need 3 n\*n matrixes)  
 $I(n) \leq n/12$ ; hence compute bound  
 $W' = (n^3)/2$  without vector ops (use 2 fma ports)  
 $W'' = (n^3)/8$  vector ops (do 4 at same time)  
 $Q' = (n^2)*3/8$  L1 & L2 (do 8 loads, 3 matrixes of size n\*n)  
 $Q'' = (n^2)*3/4$  L3 (do 4 loads, 3 matrixes of size n\*n)  
for double, multiply Q by 8 bc double needs 8 bytes

### **2.6 locality**

algorithm use data near or equal to those used recently  
more locality leads to better, more stable performance  
for example by ensuring data traversed linearly (like  $a[i][j]$ )

#### **temporal**

recently referenced items referenced again  
cycles through the same instructions

#### **spatial**

nearby addresses accessed consecutively  
instructions referenced in sequence

## **3 processor**

### **3.1 history**

memory bandwidth increased over time  
CPU speedup until 2003 (more GHz)  
since 2003 cores & vector units added

### **3.2 intel history**

fully backwards compatible (hence instructions never removed)  
end of moore's law means no longer free speedup for legacy code

#### **x86-16**

8086

#### **x86-32**

386  
486  
Pentium  
Pentium MXX (adds MXX)  
Pentium III (adds SSE; 4-way single)  
Pentium 4 (adds SSE2; 2-way double)  
Pentium 4E (adds SSE3)

#### **x86-64**

Pentium 4F  
Penryn (adds SSE4)  
Sandy Bridge (adds AVX; 8-way single, 4-way double) 32nm  
Haswell (adds AVX2) 22nm  
Skylake-V (adds AVX-512; 16-way single, 8-way double) 14nm

#### **tick-tock-opt**

shrink of same generation process technology as tick  
like sandy bridge to ivy bridge  
new generation microarchitecture as tock  
like sandy bridge to haswell  
optimization of same generation as opt  
like skylake to kaby lake (since 2016)

#### **vectorization generations**

MMX (multimedia extension)  
SSE (streaming SIMD extension)  
AVX (advanced vector extensions)

#### **super-scalar processors**

issue/execute multiple instructions in one cycle  
instructions taken from sequential instruction stream

### **3.3 memory**

#### **bottleneck**

processor performance doubles every 18 months  
but bus bandwidth only every 36 months  
for haswell, peak performance at 2 ops/cycle FMA  
needs 192 bytes/cycle, but bus bandwidth 16 bytes/cycle

#### **hierarchy**

L0 registers (words)  
L1 on-chip L1 cache SRAM (lines)  
L2 on-chip L2 cache SRAM (lines)  
L3 main memory DRAM (blocks)  
L4 local secondary storage disks (files)  
L5 remote secondary storage disks (tapes, servers)  
from smaller, faster, costly  
to larger, slower, cheaper

#### **example haswell**

16 fp registers  
L1 with 4 latency, 12 throughput (8loads, 4stores)  
L2 with 11 latency, 8 throughput  
L3 with 34 latency, 4 throughput (non-uniform)  
RAM with 125 latency, 2 throughput

### **3.4 virtual memory system**

instruction uses virtual addresses  
caches work with physical address + 12bit tag  
physical address given by lower bits of virtual address  
virtual address mapped to tag with lookup in TLB  
page size can be 4KB, 2MB, 1GB (page size step =  $2^9$ )

#### **x86 L1 cache**

64 ( $2^6$ ) block size  
64 ( $2^6$ ) sets  
8 associativity  
32 KB

#### **x86 L1 lookup**

can do concurrently with address translation  
because location inside cache given by lower 12 bits  
which are not part of page identifier (as page size  $\geq 2^{12}$ )

#### **address translation**

uses translation lookaside buffer (TLB)  
first level 128 ITLB (instructions) and 64 DTLB (data) entries  
second level (STLB) with 1536 entries  
small penalty for ITLB/DTLB miss, (very) big one for STLB miss

#### **vs L2 cache**

assume working set of  $2^{11}$  doubles with each on different page  
fits into L2 cache fully, but TLB misses

#### **MMM example**

block-column A \* whole B = block C  
if full MMM & row-major order  
only block C not contiguous (NB pages)  
if leading dimensions added  
then block B & C not contiguous (NB + K + NB pages)  
copying to continuous memory may pay off

### 3.5 instruction set architecture (ISA)

defines instruction set, registers  
implemented with microarchitecture  
which then defines caches, frequency, virtual memory  
instructions execute logic upon values in registers

#### flynn's taxonomy

single/multiple instructions (S/MI)  
per single/multiple data (S/MD)  
SISD for mainframes, old PCs with single simple core  
SIMD for GPUs, CPUs with vector instructions  
MISD for pipelines (like space shuttle)  
MIMD for scientific purposes

#### single instruction multiple data (SIMD)

for parallel computation on short int/float vectors  
allow to compute on up to 8 numbers at the same time  
useful (many applications have fine-grained parallelism)  
doable (easy to design, simply replicate units)

#### SIMD history

MMX for integers / 64bit (pentium MMX)  
SSE for 4 singles / 128bit (pentium III)  
introduces xmm0-xmm15  
SSE2 for 2 doubles / 128bit (pentium 4)  
AVX for 4 doubles / 256bit (sandy bridge)  
introduces ymm0-ymm15; three-operand instructions  
AVX2 for FMAs (haswell)  
AVX-512 for 8 doubles / 512bit (skylake-x)  
introduces zmm0-zmm31

#### registers

all instructions (including scalar) use same registers  
hence xmm0 points to the same physical address as ymm0

#### instruction names

of the form operation-size-type  
operation like add, mov, mul  
size either single (s) or packed (p)  
hence operate on single element or all  
type either single (s) or double (d)  
hence operate on 4 or 8 bytes  
like addps, movss, mulpd

#### fused multiply add (FMA)

add & multiply in same instruction ( $x + y * z$ )  
better accuracy bc only one rounding step  
natural pattern in many algorithms  
needs three operand instructions (used to be hard to build)

#### performance

latency is time to wait until result ready  
throughput is #operations at the same time  
gap is  $1/\text{throughput}$  (intel calls this "throughput")

#### example haswell

fma has latency 5, throughput 2 (gap 1/2)  
mul has latency 5, throughput 2 (gap 1/2)  
add has latency 3, throughput 1 (gap 1/1)  
same for vector instructions (except div)  
each register has space for 256 bits (4 doubles, 8 ints)

## 4 memory

### 4.1 cache

memory with short access time  
for frequently or recently used data/instructions  
supports temporal locality (by definition)  
may supports spatial locality (with lines, blocks)

#### general mechanics

data is copied/replaced in block units (intel core 64bytes)  
faster cache has subset of blocks of larger memory  
placement policy determines where block is placed  
replacement policy determines which block is replaced

#### cache organization

B for #blocks (size of single entry)  
S for #sets (number of addresses)  
E as #lines per set (associativity)  
capacity =  $S * B * E$

#### cache types

direct mapped ( $E=1$ ); each block has unique location

associative ( $E>1, S>1$ ); each block has many possible location  
fully associative ( $S=1$ ); each block to any location

#### tradeoff associativity

increase associativity if reused elements have same offset  
increase sets if then reused elements have different offsets

#### placement policy

memory locations are mapped to sets based on offset  
for constant capacity tradeoff E with S  
usually more E better (but more expensive to build)  
but for pattern 0, 1, 3, 0  $E=1, S=2$  better than  $E=2, S=1$

#### replacement policy

defines replacement if all possible locations full  
only relevant for associative caches (multiple possible locations)  
for example least recently used (LRU)

#### type of cache misses

compulsory, cold (on first access)  
capacity (working set larger than cache)  
conflict (many blocks map to same slot)

#### performance metrics

miss rate (references not in cache; misses/access)  
hit rate (references in cache; 1-miss rate)  
hit time (deliver block from cache to CPU; 3 cycles for L1)  
miss penalty (time required for a miss; 100 cycles for L2)

#### writes to memory

on cache hit write immediately (write-through)  
or defer until line replaced (write-back)  
on cache miss write directly (no-write-allocate)  
or load line and then write (write-allocate)  
intel core uses write-back and write-allocate

### 4.2 examples

#### simple

for  $E=1, B=4$  double,  $S=8$   
lower 3 bits 0 because doubles are 8 bytes  
3 bit determine sets, 2 bits for block  
data identified by 56 bit tag in single possible location

#### conflict

for  $E=1, B=2$  double,  $S=8$   
loop with  $a[j][i]$  for  $n = 16$   
hence  $a[2][0]$  conflicts with  $a[0][0]$

#### examples

for  $E=1, B=2$  double,  $S=2$   
access 0123456701234567  
8 miss, 8 hits; spatial yes, temporal no  
access 0246135702461357  
16 miss; spatial no, temporal no  
access 0123012301230123  
4 miss, 12 hits; spatial yes, temporal yes; optimal

### 4.3 traffic examples

#### vector addition ( $z = x + y$ )

memory traffic  $4n$  doubles =  $32n$  (bc measured in bytes)  
operational intensity  $1/4$  (hence memory bound)

#### naive matrix multiplication

for cache size  $y \ll n$ , block size 8  
gives  $n^2 * 9n/8$  cache misses

#### blocked matrix multiplication

for block size  $b$  such that  $3b^2 \leq \text{cache size } y$   
gives  $(n/b)^2 * 2 * nb/4$  cache misses  
 $3b^2$  bc 3 little matrix need to fit into cache

#### strided working sets problem

common in FFT, 2d transformations  
assume two loops for ( $i < n; i += t$ ) for  $t$  stride  
see how locality changes with different strides  
if stride = 1 then spatial full, temporal if  $W \leq C$   
if stride = 2 then spatial  $B/2$ , temporal if  $W \leq 2C$   
if stride = 4 then spatial  $B/4$ , temporal if  $W \leq 4C$   
(if stride not power of two not that much of an issue)  
need to separate computation in chunks  
need to copy data to better memory layout

## 4.4 memory dependencies

### read after write (RAW)

also called true dependency

$r1 = r3 + r4; r2 = 2*r1$

### write after read (WAR)

also called antidependency

$r1 = r2 + r3; r2 = r4 + r5$

dependency on  $r2$  only by name

### write after write (WAW)

also called output dependency

$r1 = r2 + r3; r2 = r4 + r5$

dependency on  $r2$  only by name

### resolve by renaming

with single static assignment (SSA) code style

automatically by compiler (but black box)

hardware can do dynamic register renaming

needs separation of architectural (assembly) / physical layer

needs more physical than architectural registers

## 5 optimizations

### 5.1 compiler optimizations

optimizations compiler is likely to do

#### mapping program to machine

register allocation

code selection & ordering

dead code elimination

eliminating minor inefficiencies

#### inline procedure calls

reduce overhead by call, bounds checking, ...

need to compile source code together

#### code motion (precomputation)

reduce computation frequency of pure computation

like moving code out of loop (loop-invariant code motion)

for example  $[i*n + j]$ ; precompute  $i*n$

#### strength reduction

replace costly operation with cheaper one

like shift/add instead of multiply by 2

for example  $[i*n + j]$

transform  $i*n$  to addition in outer loop

#### share common subexpression (if recognized)

reuse portion of expressions

for example  $[(i+1*n) + j] [(i*n) + j+1]$

replace with  $[c + n]$  and  $[c + 1]$  for  $c = i*n + j$

### 5.2 compiler optimization blockers

even unrealistic scenarios must not change

analysis constrained (only static procedure-wide info)

can't test different choices (for ILP, locality)

hence algorithmic restructuring hard

#### procedure calls

treated as black box with potential side effects

hence can not simply reorder, precompute

compiler may detect pure (built-in) functions

#### memory aliasing

overlapping memory prevents reuse optimizations

like reducing loads/stores, precomputation

compiler may add case distinction for aliasing ( $a+n < b \parallel b+n < a$ )

and then use optimized code when no memory alias

### 5.3 processor optimizations

#### out-of order-execution

dynamically reorders instructions

to resolve dependencies early & use free ports

#### register renaming

resolve WAW, WAR dependencies

#### branch prediction

execute instructions before branch result known

result instantaneously ready if branch result predicted correctly

## 5.4 manual optimizations

### algorithm

restructure for ILP, locality

use simple data representations; best are arrays

watch inner-most loop closely for overhead

unroll loops to enable other optimizations

try out different algorithms

### optimize procedure calls

reduce overhead when calling procedures

precompute manually (taking out of loop & similar)

provide source code instead of linking (to enable inlining)

### scalar replacements (avoid aliasing)

copy reused elements in local variables

structure algorithm in load, compute, store

### unroll loops

unroll loop by some factor

do scalar replacement to save on loads

measure outcome and autotune unrolling factor

### generate code

capture algorithm with domain specific language (DSL)

decide on choices with bruteforce/machine learning

generate easily transformable code to profit from compiler

### tell compiler no aliasing

use compiler flag `-fno-alias`

use `#pragma ivdep`

pass argument with restrict keyword

## 5.5 loop unrolling

more work inside single loop iteration

faster iterator incrementation

but just unrolling usually no advantage

need to transform (reassociation, dependency reduction)

### example

create two accumulators  $i0, i1$

sum up  $i$  and  $i+1$  into respective accumulator

sum up  $i0 + i1$  in the end

### compilers

transformations may break associativity

hence compiler not allowed to do so

many different forms of how to unroll

hence compiler search space too big

### generalization

use  $K$  accumulators and  $L$  unrolls (for  $K$  divides  $L$ ;  $L \geq K$ )

choose  $K/L$  combination with best performance

but processor-specific hence often unportable

called autotuning, processor-tuning

### approaches

try out  $K, L$  then measure and choose best

but large overhead for small inputs

pick best  $K, L$  using latency/throughput model

but models may not work in practice

## 5.6 loop unrolling examples

based on sum of array

### trivial

for (entry in array)  $\text{sum} += \text{entry}$ ;

each computation pays full latency

because always dependency on operation before

### loop unrolling

for (entry, next\_entry in array)  $\text{sum} = (\text{sum} + \text{entry}) + \text{next\_entry}$ ;

### loop unroll, separate accumulator

for (entry, next\_entry in array)  $\text{sum1} += \text{entry}$ ;  $\text{sum2} += \text{next\_entry}$

now dependencies reduced by half

but for floats behaviour changes

for integer fine bc. overflow preserves associativity

### loop unroll for mul

$\text{ceil}(5 \text{ lat}/0.5 \text{ gap})$  operations need independence

hence choose  $L=K=10$

### loop unroll for add

$\text{ceil}(3 \text{ lat}/1 \text{ gap})$  operations need independence

hence choose  $L=K=3$

but best performance measured at L=K=8

### loop unroll for add using fma

ceil(5 lat/0.5 gap) operations need independence  
another 2x speedup bc one additional port used

## 6 vectorization

use SIMD instructions to parallelize computations

### 6.1 implementation options

vectorized library (but seldom available)  
compiler flags (but seldom works)  
intrinsics (but limited optimisations possible)  
assembly (but no help anymore from compiler)

### 6.2 data layout

aliased data can not be vectorized (result could be different)  
unaligned data has slower loading performance

#### compiler decision tree

if a,b aliased then unvectorized code  
else if a,b aligned then load aligned  
else use unaligned loads  
or execute first entries unvectorized ("peeling")  
impractical to do on large code bases

#### detect vectorization

annotation flags of compiler  
generate vectorization report

#### improve

prevent aliasing with flags, pragmas or restrict keyword  
avoid branches like switch, goto, return (straight-line code)  
terminate loop by constant, linear function, loop invariant  
inner loop with consecutive, dependencyfree access  
access elements directly (with loop variable, index)  
use arrays, not structs / pointers  
use #pragma vector aligned or \_\_assume\_aligned(vec, 32)

### 6.3 intrinsics

C functions translating to specific assembly  
intel intrinsics guide for detailed specification  
same instruction in parallel easy & cheap (like add)  
crossing lines difficult & expensive (like shuffle)

#### as an abstraction layer

need to explicitly load/store operated variables  
need to have data aligned to 256bit (32bytes)  
parameter variables not overwritten by operation  
(not even if assembly would, like fma)

#### flavors

SSE with 4-way floats, 2-way doubles  
%xmm0-%xmm15 128bits registers  
AVX with 8-way floats, 4-way doubles  
%ymm0-%ymm15 256bits registers  
AVX-512 with 16-way floats, 8-way doubles  
%zmm0-%zmm31 512bits registers  
mixing flavors sets costly  
bc need to save/restore upper bits

#### function naming

.\_<prefix>.\_<operation>.\_<suffix>(<data>)  
prefix like \_\_mm, \_\_mm256, \_\_mm512  
operation like load, add, store  
suffix like size-type of instruction  
data like \_\_m128, \_\_m256, \_\_m512 or naive types  
like \_\_mm256\_load\_pd(4.0,3.0,2.0,1.0)

#### types

native (one-to-one to assembly) like \_\_mm256\_load\_pd()  
multi (maps to several) like \_\_mm256\_set\_pd()  
macros (simplify code) like \_\_MM\_SHUFFLE()

#### data types

\_\_m256, for float[4]  
\_\_m256d for double[2]  
\_\_m256i for char[16], short int[8], int[4] or uint64\_t[2]

#### loads

load\_pd(pointer) → pointers must always be 32bit aligned

loadu\_pd(unaligned pointer; used to be more expensive)  
loadh\_pd(\_\_m128d t, unaligned pointer p) → [t1, p]  
maskload\_pd(pointer, \_\_m256i mask with 0x0 and 0xFFFF)  
broadcast\_sd(one double repeated to fill line)  
boradcast\_pd(two doubles repeated; for complex numbers)  
i64gather\_pd(pointer, \_\_m256i offsets, scale; 8 for doubles)

#### sets

set1\_sd(3.0) → [3.0, 3.0, 3.0, 3.0]  
set\_pd(4.0, 3.0, 2.0, 1.0) → [4.0, 3.0, 2.0, 1.0]  
setr\_pd(4.0, 3.0, 2.0, 1.0) → [1.0, 2.0, 3.0, 4.0]  
set\_sd(4.0) → [0.0, 0.0, 0.0, 4.0]  
setzero\_pd() → [0.0, 0.0, 0.0, 0.0]  
set\_m128d(128bit low, 128bit high) → [low, high]

#### arithmetic

add, subb, mul, div, sqrt  
max, min, ceil, floor, round  
addsub(a, b) → [a1+b1, a2-b2, a3+b3, a4-b4]  
hadd(a, b) → [a1+a2, b1+b2, a3+a4, b3+b4] (also hsub)  
fmadd(a, b, c) → [a1\*b1 + c1, ...] (also fmsub, fmaddsub)  
cmp(a, b) → [if a1=b1 then 0xFF else 0x00, ...]  
dp(a, b, int mask\_sum |mask\_result) for dotproduct

#### comparison

cmp\_pd(a, b, mode) for mode = <operation>.\_<ordering>  
operation is EQ (=), GE (≥), GT (>), LE (≤), LT (<)  
to negate, prefix operation with N  
ordering is OQ (ordered; includes NaN), UQ (unordered; ignores NaN)  
like cmp(a,b, \_\_CMP\_EQ\_OQ), cmp(a,b, \_\_CMP\_NEQ\_UQ)  
special operations are \_\_CMP\_TRUE\_UQ, \_\_CMP\_ORD\_Q (and vice-versa)  
result is 0xFFFF when fulfilled, and 0x0 otherwise

#### conversions

epi32 for 32bit integers  
cvt\_epi32\_pd(\_\_m128i a) → \_\_m256d (4 ints to doubles)  
cvt\_epi32\_ps(\_\_m256i a) → \_\_m256 (8 ints to floats)  
cvt\_pd\_epi32(\_\_m256d a) → \_\_m128i (4 doubles to ints)  
cvt\_pd\_epi32(\_\_m256 a) → \_\_m256i (8 floats to ints)  
cvt\_ps\_pd(\_\_m128 a) → \_\_m256d (4 floats to doubles)  
cvt\_pd\_ps(\_\_m256d a) → \_\_m128 (4 doubles to floats)  
cvt\_tpd\_epi32(\_\_m256d a) → \_\_m128i (4 doubles to ints)  
cvt\_sd\_f64(\_\_m256d a) → double (first double extract)  
cvt\_ss\_f32(\_\_m256 a) → float (first float extract)

#### shuffles

unpackhi\_pd(a, b) → [a2, b2, a4, b4]  
unpacklo\_pd(a, b) → [a1, b1, a3, b3]  
movemask\_pd(a) → 4-bit int; bit[i] = MSB of a[i]  
movedup\_pd(a) → [a1, a1, a3, a3]  
blend\_pd(a, b, mask) → [a1 or b1, a2 or b2, ...]  
blendv\_pd(a, b, (variable) mask) → [a1 or b1, a2 or b2, ...]  
insertf128\_pd(a, b, 0 or 1) → [b1, b2, a3, a4] or [a1, a2, b1, b2]  
extractf128\_pd(a, 0 or 1) → [a1, a2, 0, 0] or [a3, a4, 0, 0]  
shuffle\_pd(a, b, mask) → [a1 or a2, b1 or b2, a3 or a4, b3 or b4]  
permute\_pd(a, mask) → [a1 or a2, a1 or a2, a3 or a4, a3 or a4]  
permute4x64\_pd(a, mask) → [any a\*, any a\*, ...] (but 128bit cross)  
permute2f128\_pd(a, b, mask) → [a1 or a2 or b1 or b2 or 0, ...]

#### stores

store\_pd(pointer); pointer must be 32bit aligned  
storeu\_pd(unaligned pointer); used to be more expensive  
maskstore\_pd(pointer, \_\_m256i mask with 0x0 and 0xFFFF)  
storeu2\_m128d(low bits pointer, high bits pointer)  
stream\_pd(pointer) (no caching)

### 6.4 intrinsic examples

#### add own index to array x

(for i <x.size(); i+=4)  
vec = load\_pd(x+i)  
index = set\_pd(i+3, i+2, i+1, i)  
vec = add\_pd(vec, index)  
store\_pd(x+i, vec)

#### add own index to array x

index = set\_pd(3,2,1,0)  
increment = set1\_pd(4)  
(for i <x.size(); i+=4)  
vec = load\_pd(x+i)  
vec = add\_pd(vec, index)  
store\_pd(x+i, vec)  
index = add\_pd(index, increment)

#### add 1 if negative, else add -1

```

ones = set1_pd(1)
negatives = set1_pd(-1)
comparator = set1_pd(0)
(for i <size; i+=4)
values = load_pd(x+i)
comparison = cmp_pd(values, comparator, _CMP_GE_OQ)
addition = blendv_pd(ones, negatives, comparison)
result = add_pd(values, addition)
store_pd(x+i, result)

```

#### load 4 from arbitrary memory

```

v1 = loadu_pd(p1), v2 = loadu_pd(p2)
v3 = loadu_pd(p1-2), v2 = loadu_pd(p2-2)
v1v2 = unpacklo_pd(v1, v2)
v3v4 = unpacklo_pd(v3, v4)
v = blend_pd(v1v2, v3v4, 0b0011)

```

#### load 4 from arbitrary memory with $set_{pd}$

```

v = set_pd(p4, p3, p2, p1)
reverse engineered the same as
v1 = _mm_load_sd(p1), v3 = _mm_load_sd(p3)
v1v2 = _mm_loadh_pd(v1, p2), v3v4 = _mm_loadh_pd(v3, p4)
v1v2d = _mm256_castpd128_pd256(v1v2)
v = _mm256_insertf128_pd(v1v2d, v3v4, 1)

```

## 7 roofline plot

way to plot performance vs operational intensity  
visualizes how well the algorithm is optimized to machine  
x log axis is  $I(n) = W(n) / Q(n)$   
y log axis is  $P(n) = W(n) / T(n)$

### 7.1 data

work  $W(n)$  in [flops]  
data movement  $Q(n)$  in [bytes]  
runtime  $T(n)$  in [cycles]

### 7.2 roofs

#### peak performance (pi)

in [flops/cycle], per core  
can never be faster then peak performance  
draw horizontal line at  $y = \text{peak performance}$

#### memory bandwidth (beta)

in [bytes/cycle], per memory layer  
use stream benchmark to measure (could be conservative)  
measured values up to 60% of theoretical possible from manual  
 $B \geq Q / T = P/I \Rightarrow \log(P) \leq \log(I) + \log(B)$   
 $y=x$  line with offset to cut peak performance root at  $\text{pi}/\text{beta}$

#### interpretation

left of  $\text{pi}/\text{beta}$  is memory bound, right is compute bound

### 7.3 observations

#### introducing SIMD

increases pi roof  $\Rightarrow$  more memory bound  
bc can execute many operations in parallel

#### different algorithmic operation mix

lowers pi roof  $\Rightarrow$  more compute bound  
bc different available units/ports

#### introducing parallelism

increases performance, I ideally stays the same  
because same work/data movement should stay same

#### warmup cache

I gets huge for small problem sizes  
for large no effect bc problem does fit in cache

### 7.4 examples

get lower bound of  $Q$  with cold misses

#### $y = ax+y$ (in doubles)

$W = 2n$   
 $Q = 16n$  (read) +  $8n$  (writes) =  $24n$   
 $I = 1/12$   
in plot, hits capacity line  
for all inputs sizes at the same point  
for parallel, hits increased capacity line

#### $y = Ax + y$

$W = 2n^2$   
 $Q \geq 8n^2$  (+24n)  
 $I = 1/4$   
in plot, nearly hits capacity line  
for larger inputs sizes gets closer to memory line (vertical)  
for parallel, hits increased capacity line

#### $C = AB + C$

$W = 2n^3$   
 $Q \geq 32n^2$  ( $\geq$  bc capacity misses for large matrixes)  
 $I \leq n/16$   
in plot, nearly hits peak performance  
for small inputs  $I$  increases (as long as cache available)  
for large inputs,  $I$  only slightly smaller (bc of blocked implementation)  
for parallel, nearly hits increased peak performance  
for parallel, gets closer to hit capacity line  
good algorithms are compute bound  
trivial triple loop is memory bound

## 8 sparse linear algebra

linear algebra usually memory bound  
hence for sparse matrixes, compression pays off  
for BLAS 1, only 18% performance in 8mb working set

### 8.1 sparse matrix vector multiplication (MVM)

core building block of sparse linear algebra  
 $y += A * x$ ; executed many times for different  $x$ 's

#### operational intensity

for  $K$  nonzero entries  
 $W(n) = 2K$  (one add, one mul per entry)  
 $Q(n) = K + 3n$  (load nonzeros; load  $y, x$ ; store  $y$ )  
hence upper bound  $\leq 1/4$  for doubles ( $8 * Q(n)$ )

#### trivial storage format

many zeros stored; unnecessary operations & movements  
operational intensity very low  
hence need better storage formats

### 8.2 compressed sparse row format (CSR)

NM matrix  $A$ ,  $K$  non-zero entries  
values with non-zero entries in row-major ( $K$  length)  
 $col_{idx}$  with column of respective value ( $K$  length)  
row\_start with index where new row starts ( $m+1$  length)

#### storage

$2K$  (values & column array) +  $m+1$  (row\_start array)

#### code

loop over  $m$  rows  
go to next row when  $col_{idx}[i] \geq row\_start[j]$

#### example

$[a, 0, b; 0, 0, 0; 0, c, 0]$   
values  $[a, b, c]$   
 $col_{idx} [0, 2, 1]$   
row\_start  $[0, 2, 2, 3]$

#### advantages

only nonzero values stored  
all arrays accessed consecutively (good spatial locality)  
good temporal locality with  $y$

#### disadvantages

insertion into  $A$  costly  
bad spatial & temporal locality with  $x$  of  $y = A * x$   
because columns (hence matching entry of  $x$ ) unordered

#### comparison to dense MVM

$2x$  slower compared to trivial loop  
performance is input dependent  
blocking requires change of data structure

### 8.3 blocked CSR (BCSR)

block for registers to reduce traffic with  $x$   
NM matrix  $A$ ,  $K$  non-zero entries,  $r * c$  blocksize  
store whole blocks including zero-values  
 $b\_values$  with all blocks which have non-zero entries  
 $b\_col\_idx$  with block-based column of block  
 $b\_row\_start$  with index where block-based row starts

### storage space

rcK (blocks) + K (values) + m/r+1 (row\_start)

### code

loop over m/r+1 rows  
dense micro MMM in inner loop

### example

```
[a,0,b,0; 0,0,0,0; 0,0,0,0; 0,c,0,0]
for r=c=2
values [a,0,0,c,b,0,0,0,0]
col_idx [0,1]
row_start [0,1,2]
```

### advantages;

temporal locality of x and y  
less index storage

### disadvantages

need to store some 0s (storage, computational overhead)  
relevant bc problem already memory bound

## 8.4 choose CSR / BCSR

performance hard to predict, machine dependent  
BCSR often better in well-known sparse matrix set  
but CSR also has winning scenarios

### by extensive search

transformation cost CSR → BCSR equals 10 SMVM  
hence extensive BCSR block size search too slow

### by estimating gain/loss

estimate gain with dense performance of BCSR / CSR  
is machine dependent, hence could do at compile time  
estimate loss with matrix values of BCSR / CSR  
is data dependent, hence likely need to do at run time  
choose best gain / loss ratio

## 8.5 principles

### optimize memory hierarchy

blocking for registers (same as ATLAS)  
change data structure of A  
optimizations are input dependent

### fast basic blocks (micro MVM)

unrolling / scalar replacement

### search for CSR, BCSR & block size

use performance model to choose best variant  
(versus measuring runtime like in ATLAS)

## 8.6 other ideas

### cache blocking

divide sparse matrix into smaller sparse matrixes  
only useful for very large matrixes

### value compression

only store unique values  
replace values array with index to its unique value

### pattern-based compression

extend BCSR with bit patterns defining non-zero entries  
choose different kernel for each bit pattern  
in values, no longer need to store 0s now  
like 0306 → bit pattern 0101; values 142

### multiple inputs scenarios

block across multiple MVM runs  
hence do A \*xyz in single run

## 9 compilers

### 9.1 gcc

no optimizations on as default

### flags

-O2, -O3 (group of optimization flags)  
-march=native (architecture)  
-mAVX (use avx instructions)  
-m64 (64 bit architecture)

### matrix multiplication example

1300 cycles at -O0

150 cycles at -O3 -m64 -march

## 9.2 exercise tools

### numbers

$2^{10} = 1\text{kb}$   
 $2^{20} = 1\text{mb}$

### bounds

runtime like 3N cycles  
performance like 2flops/cycle  
performance = #flops / #cycles

### operational intensity

for W, include reads + writes & multiply by data size  
 $W(n)$  of  $A*y = 2(n^2)$   
 $W(n)$  of  $A*B + C = 2(n^3)$   
 $I < \pi/\text{bandwidth}$  for memory bound  
 $W/\pi = Q/\text{beta}$  for tipping point of compute/memory bound  
 $Q(n)$  of  $Y = X + Z = 3(x^2)$  bc have to include load + store  
 $I(n)$  accurate if all fits in cache  
write-through cache does not allocate result

### cache

offset of storage relevant for set placement  
miss rate = miss / total \*100  
spatial if neighbours accessed  
temporal if same element used multiple times  
levels are L1, L2, ..., DRAM

### roofline

intersection point x-axis at peak / bandwidth in bytes  
performance within memory bounds  $\leq B * I$

### CSP

#row\_start = #rows + 1  
last number row\_start = #non-null elements